


# COMP 4735: Operating Systems Concepts

## Lecture 2: What is an Operating System?



Rob Neilson

[rnelson@bcit.ca](mailto:rnelson@bcit.ca)


# Reading

- The following sections should be read before next Monday
  - Textbook Sections: 1.1, 1.2, 1.3 (Theory Part)  
10.1, 11.1 (Case Study Part)
- Yes, there will be a quiz next Monday in lecture
  - A sample quiz will be posted on webct on Friday
  - The sample quiz will be reviewed next Monday in lecture, before the real quiz
  - This quiz (Quiz 1) covers material from all 5 of the sections listed above
    - You might do well to focus on the 'key concepts' presented in lecture

# Typical Computer

- A typical modern computer consists of (minimally):
  - One or more processors
  - Main memory (volatile storage)
  - Disks (permanent storage)
  - Various input/output devices
- Managing all these components requires a layer of software ...  
... the **operating system** ...

# ENIAC (1946): A Typical Computer?



- thirty separate units
- 19,000 vacuum tubes
- 1,500 relays
- hundreds of thousands of resistors, capacitors, and inductors
- required 200 kilowatts of electrical power
- needed forced-air cooling
- weight: 30 tons

- ENIAC could:
  - discriminate the sign of a number
  - add, subtract, multiply, divide
  - compare quantities for equality
  - extract square roots
  - ENIAC stored a maximum of twenty 10-digit decimal numbers.

# ENIAC Operating System



Which OS?  
... here he is!

- ENIAC did not have an Operating System.
- Programs were 'loaded' by plugging tubes and relays together using cables



# Life kinda sucked for ENIAC programmers ...

- ENIAC could only run one program at a time
- it had no Operating System per se
- the programmers had to move every byte or bit manually in and out of registers, accumulators etc as needed
- they plugged cables and flipped switches to set an initial state, and then turned it on
- there was no memory

“The ENIAC performed arithmetic and transfer operations simultaneously, but *concurrent operation caused programming difficulties.*”

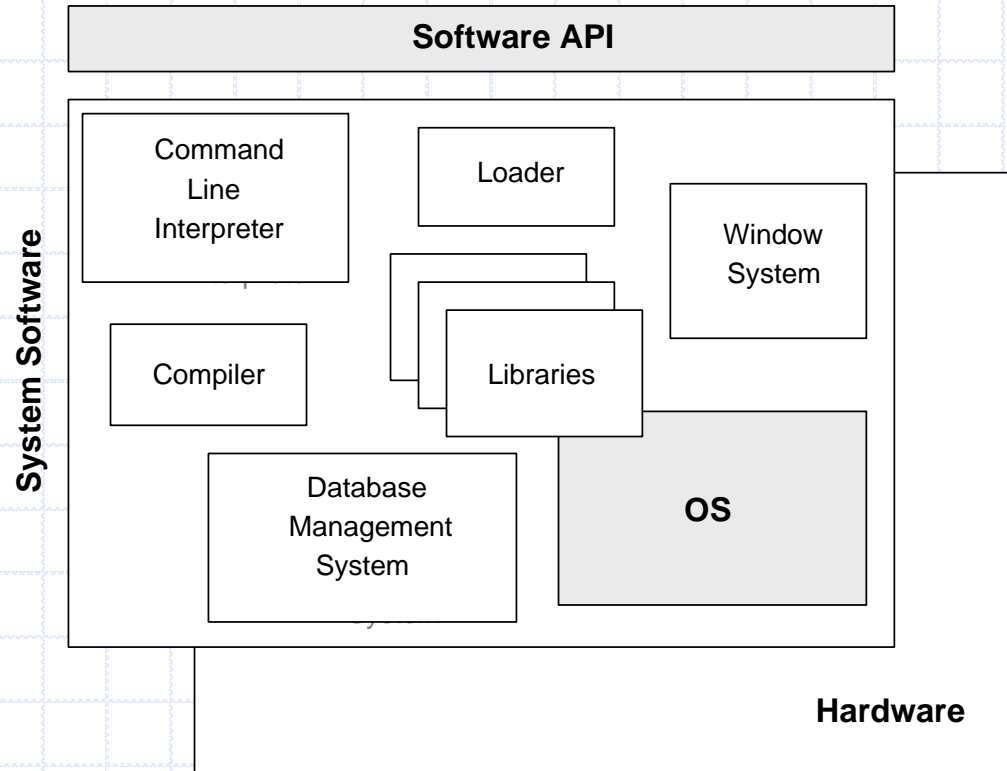
- Martin H. Weik, 1961

Ordnance Ballistic Research Laboratories

# What was the problem ...

- The thing needed a better way to control it  
(actually, it needed a lot of things, but from a programmers perspective an OS would have helped)
- So what would an OS have done for them?
  1. provide an abstract interface to all that hardware (***resource abstraction***)
  2. allow resources to be shared (***resource sharing, or multiplexing***)
    - they got into trouble when I/O and processing (CPU) had to be done concurrently
    - what they needed was a ***super-smart dude*** who knew which wires could be plugged into which hole such that concurrent operation worked
      - this dude would be some kinda “supervisor”, or, in a modern computer, the dude is the OS

# The Operating System and System Software



- In our study of operating systems, we care about the two grey boxes ...
  - the Software API
    - provides an programming interface to the OS
  - the OS itself
    - abstracts and controls the hardware and resources

- The main idea is that all this stuff is the plumbing that you need to write programs – applications – to make the computer do something.



# Key Concept: Resource Abstraction

- One fundamental thing any OS must do is to provide resource abstraction
- What exactly do we mean when we say “resource abstraction”?
- Well, the resources are things like disk space, cpu time, registers, etc – but we want to perform higher level ops like “write Hello World to the console”

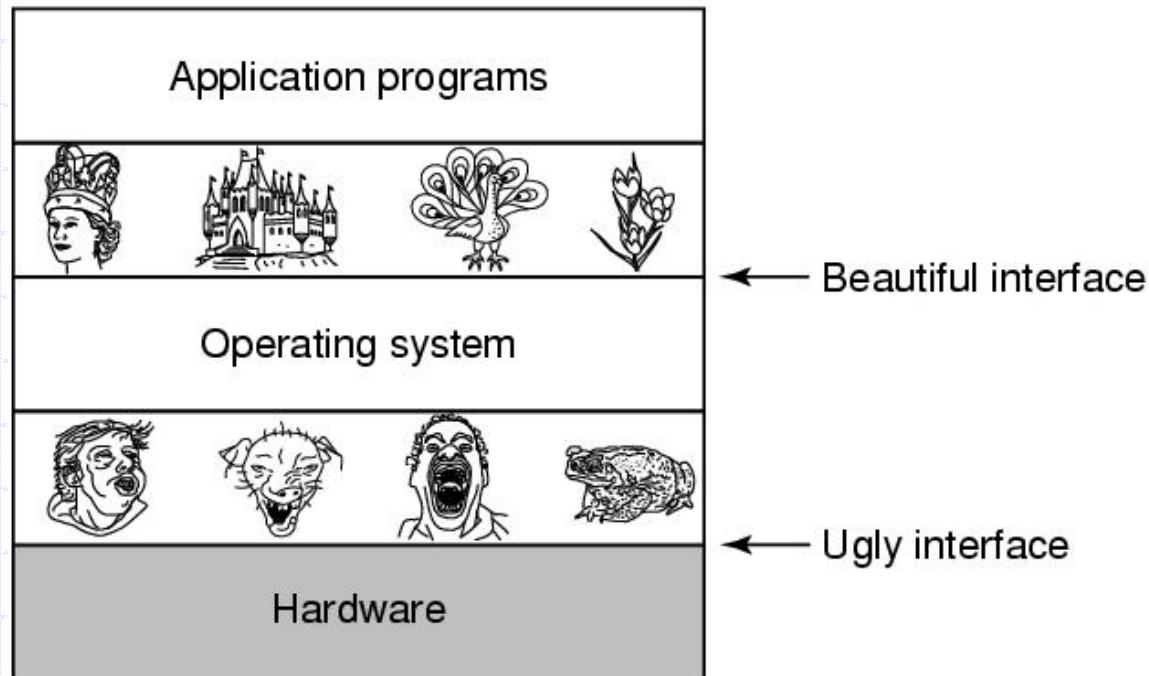
```
load(...);  
seek(...);  
out(...);
```

```
void write() {  
    load(...);  
    seek(...);  
    out(...);  
}
```

```
int fprintf(...) {  
    ...  
    write(...);  
    ...  
}
```

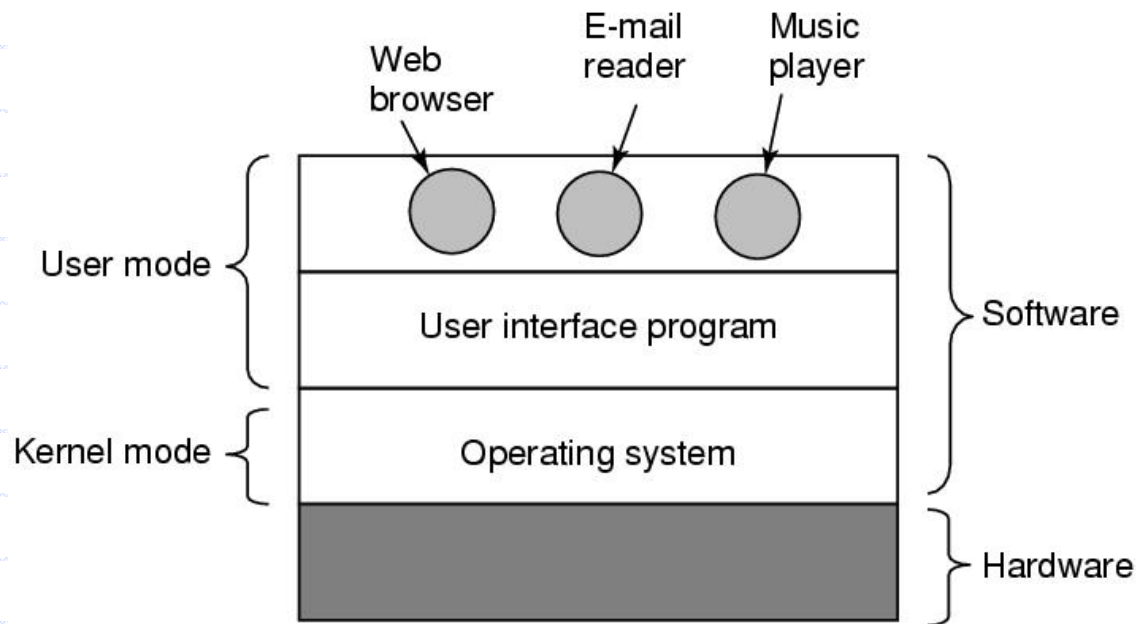
- The diagrams above show progressive levels of abstraction ...
  - on the left we have basic disk drive control cmds
  - on the right we have higher-level programming API calls (stdio.h)
- Each level hides details from lower levels, but ...
  - ... the trade-off is a loss in flexibility
- *Can you guess what flexibility might have been lost in the example above?*

# OS vs. Hardware Interfaces



- Hardware interfaces are complex and ugly
- The OS needs to hide the device specific details from the programmer.

# Where the OS fits in ...



- A modern OS uses the concepts of *user mode* and *kernel mode*
- User mode software is stuff that the user or programmer might want to control
- Kernel mode stuff is critical for the operation of the computer .... we don't let the users play with it because they may cause the computer to crash

# Key Concept: Resource Sharing

- The OS is basically a “resource manager”. This is its main purpose.
  - resources are things like memory, disk space, CPU
- The OS is required to:
  - Allow multiple programs to run at the same time
  - Manage and protect memory, I/O devices, and other resources
  - Includes multiplexing (sharing) resources in two different ways:
    - In time
    - In space

# The OS as a Resource Manager

- Let's consider a real world example...
  - The following pages are from a book titled "DMS SuperNode System Description"
    - This is the OS used in older Nortel digital phone switches
  - These pages are the contents of one chapter that describes the OS nucleus – or kernel.
  - Notice that the OS kernel is basically one big resource manager
  - Look at the types of things this OS manages ...
    - Memory, processes, timers, queues, messages, etc



<Insert DNS PDF files here>

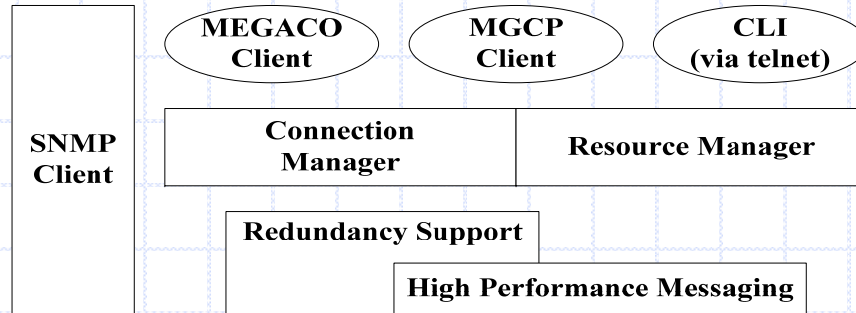


# Another example of abstraction

- consider an embedded OS
  - has no user interface component
  - just has an API
- the “target” board contains the embedded processor and a bunch of other resources, in this case
  - switching channels
  - DSPs
    - to implement vocoders, codecs, protocols etc
  - the basic services on the board are accessed by a very basic API that uses a 1:1 mapping between C function calls and processor opcodes
- the OS needs to provide an abstraction
  - what we care about is making phone calls, so how about ...


# Layers of abstraction in a sample API ...

## Connection Layer




The highest level of abstraction implements and manages the concept of 'connections' between phones

## Service Layer



This level combines system calls to chips to provide higher level services, such as send/receive of messages

## DSP Core



These APIs abstract individual physical chips on the board

Bottom layer just communicates to the board

**HARDWARE**

# OS Responsibilities...

The OS has a contract. It has an obligation to the rest of the system. This obligation includes:

- The OS has to make sure that programs *only use resources when they are available*
- The OS has to make sure that programs *only access resources that they are allowed to*
- The OS has to make sure that many programs can *share access to system resources* .... ie ... OS facilitates '*Multiprocessing*'
- The OS has to be *trusted* ... ie ... the OS has to work exactly as specified. This part of the system is *mission critical and cannot fail*

# A Short History of Operating Systems

## Generations:

(1945–55)	Vacuum Tubes
(1955–65)	Transistors and Batch Systems
(1965–1980)	ICs and Multiprogramming
(1980–Present)	Personal Computers

# Batch Systems




IBM 370 (1970)



- OS processes “jobs”
- “jobs” are non-interactive
- “jobs” are controlled by a sequence of input commands

# Key Concept: Batch Processing



- (a) Programmers bring cards to 1401.
- (b) 401 reads batch of jobs onto tape.
- (c) Operator carries input tape to 7094.
- (d) 7094 does computing.
- (e) Operator carries output tape to 1401.
- (f) 1401 prints output.



# A Batch Job (Card Stack for a Job)

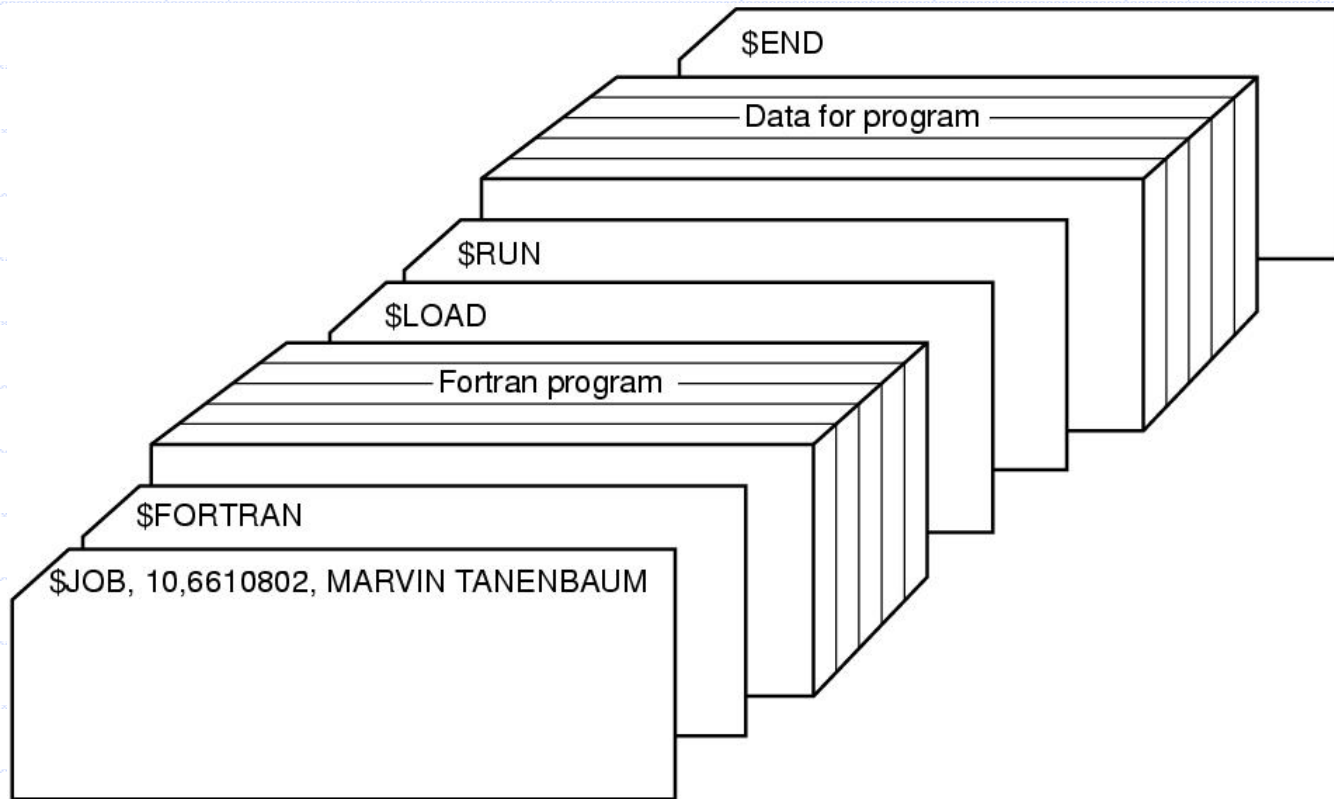


Figure 1-4. Structure of a typical batch job.


# Batch Processing (Sample JCL)

```
//SIMOJOB1 JOB (ACCTINFO),CLASS=A,MSGCLASS=C,NOTIFY=USERID
//*
//TEMPLIB1 DD DISP=(NEW,CATLG),DSN=&DSNAME,
//          STORCLAS=MFI,
//          SPACE=(TRK,(45,15,50)),
//          DCB=(RECFM=FB,LRECL=80,BLKSIZE=800,DSORG=PO)
//*
//JOBLIB DD DSN=SIMOTIME.DEVL.LOADLIB1,VOL=DCB083
//*
//STEP0100 EXEC PGM=PROGRAM1
//*
//STEP0200 EXEC PGM=PROGRAM2
//*
//STEP0300 EXEC PGM=PROGRAM3
//STEPLIB DD DSN=SIMOTIME.DEVL.TEMPLIB1,DISP=SHR
//*
//STEP0400 EXEC PGM=PROGRAM4
```

# Key Concept: Multiprocessing

- ***Multiprocessing*** is the term given to the sharing of a computer and its resources by more than one user, or more than one job
- Operating systems that supported batch processing were the first to incorporate multiprocessing
- The idea was to make efficient use of CPU time when jobs were waiting for I/O devices or for operator intervention

# Multiprocessing of Batch Jobs



- Memory is divided into fixed partitions
- Each job is loaded into a partition
- OS runs on job at a time until the job blocks (eg: waiting for a tape to be mounted etc)
- When a job is blocked the CPU will switch to a different partition

Note:

- resources such as tape drives, printers, volumes (storage disks) are allocated to a job and held until the job finishes
- only one job can use a resource at a time

# Other Batch Systems



Dennis Richie and Ken Thompson's design of the first version of UNIX on a PDP-7 marked the beginning of a long relationship between DIGITAL and UNIX that continues to the present. In the early 1970s UNIX was ported to a PDP-11 and in 1983, DIGITAL developed ULTRIX, its own version of the popular operating system. Here, Richie and Thompson work at a PDP-11.

# Unix batch man page

## NAME

batch - queue, examine or delete jobs for later execution

## SYNOPSIS

**batch** [-V] [-q *queue*] [-f *file*] [-mv] [**TIME**]

## DESCRIPTION

**batch** reads commands from standard input or a specified file which are to be executed at a later time, using the shell set by the user's environment variable **SHELL**, the user's login shell, or ultimately **/bin/sh**.



# Key Concept: Timesharing Systems



# Timesharing

- Uses multiprogramming to enable multiple 'simultaneous' users
- Support interactive computing model (illusion of multiple consoles)
- Different scheduling & memory allocation strategies than batch
- Considerable attention to resource isolation (security & protection)
- The goal was to optimize to optimize response time for users

# PCs and Workstations

- Fourth generation Operating Systems are what we have today with XP, Vista, Linux, MacOS etc
- These OSs are designed primarily for single users (although they support multiple users)
- Make use of inexpensive hardware resources
- Initially these were 'bare bones' systems (eg: CP/M) to run programs on
- Now they are sophisticated and more complex than any of their predecessors

*We will discuss these in more detail when we look at the Linux and Vista case studies*

# Summary of Key Concepts

- OS manages the physical resources of the computer
- OS provides resource abstraction to better let us use the computer as a tool (ie: to perform tasks <like gaming> that we want to perform)
- OS makes sure that resources are shared safely and fairly between processes
- OS enables multiprogramming
- OS needs an interface for programmers, or it is useless
- OS typically (but not always) needs a user interface
- Batch systems work in an “offline” mode
- Timesharing systems work in an “online” mode
- Modern OS’s support both batch and timesharing operations



# The End