

COMP 4735: Operating Systems

Lecture 7.1: Threads Part 2



Rob Neilson

rneilson@bcit.ca

Reading

- The following sections should be read before next Monday
 - it would also help to read this before your lab this week
 - Textbook Sections: 2.2 (Theory Part)
10.3.3 (Case Study Part)
- Yes, there will be a quiz next Monday in lecture on the above sections
 - A sample quiz will be posted on webct on Friday
 - This quiz (Quiz 4) covers material from all of the sections listed above

Agenda

Key concepts for this lesson:

- User Space Threads
- Kernel Space Threads
- Threads in Linux

Question 5

5. Assume that you compile and run the program shown below on a multi-threaded OS. How many threads are created?

- a) zero
- b) one
- c) two
- d) three
- e) four

```
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>

void *bye (void *id) {
    printf ("I am about to die ... %d\n", id);
    pthread_exit(0);
}

int main (int argc, char *argv[])
{
    int i=1;
    int rc;
    pthread_t threads[3];
    rc=pthread_create(&threads[i], 0, bye, (void *)i);
    exit(0);
}
```


Question 6

6. Which of the following Pthread library calls would you make if you want your program to block and wait for a specific thread to finish executing?

- a) `Pthread_block()`
- b) `Pthread_exit()`
- c) `Pthread_join()`
- d) `Pthread_wait()`
- e) `Pthread_yield()`
- f) none of the above


User Space Threads

- threads are implemented in a package (library)
- a multi-threaded program must include the library; (eg: Pthreads)
- the code in the library implements everything that is needed to support multi-threading
 - thread table
 - thread scheduler
 - thread state machine
 - all of this information is managed by a thread **run-time system**
- **the threads all execute within a process**
- **the CPU is completely unaware of the threads**
- **the kernel schedules processes not threads**



User Space Threads - Management

- threads are managed by the run-time system
- when a thread wants to create another, a library call is made
- this library call is local and is handled within the process
- the thread tables are managed within the process - without ever trapping to the kernel



User Space Threads - Benefits

- thread switching is very **fast** as there is **no need to trap to the kernel**
- can run on **legacy systems** that do not support kernel threads
- **scale better** (ie: do not affect system performance as much as kernel threads when the number of threads gets large)
 - ... because they **don't require kernel memory** for tables etc

enable lightest-weight, fastest method of switching threads of execution

User Space Threads - Drawbacks

- blocking system calls cause *entire process to block*
 - this means that *all threads in a process are blocked* every time any thread wants to do IO or has a page fault
 - unfortunately the biggest benefit of multi-threading is for programs that do lots of IO ... so we have a *paradox*
- threads have to *voluntarily give up the CPU*
 - if a thread hogs the CPU, other threads in the process *will not run*
 - the thread scheduler has *no way of pre-empting a thread*
 - the *clock* in the kernel can *pre-empt the entire process*, but not a specific thread
 - the threads have to call *thread_yield()* to voluntarily give up the CPU

Question 7


7. Assume you have a process that is running a multi-threaded program on a system that uses user-space threads. Which one of the following events will cause all the threads in the process to block?
- a) the garbage collector thread starts executing
 - b) a thread calls `thread_create()` to create a new thread
 - c) a thread calls `thread_exit` to terminate itself
 - d) a thread calls `read()` to load a record from a file
 - e) a thread is scheduled for execution
 - f) none of the above

Question 8

8. Assume that a new thread is created on a system that runs a user space threads package. The thread is put in the READY state. Where will this status information (READY) be stored?
- a) in a process table in the process
 - b) in a process table in the thread
 - c) in a thread table in the kernel
 - d) in a thread table in the process
 - e) none of the above

Kernel-Space Threads

- user space threads were implemented in a library, and the OS knew nothing about them
- **kernel threads are implemented inside the kernel itself**
- the OS designers have to create a process table extension (ie: **thread table**), and **change the kernel scheduling and context switching components** to operate on threads
- **the kernel schedules threads - not processes**
- **the CPU manages the threads, including:**
 - **creation**
 - **termination**
 - **scheduling**
- **a program must trap to the kernel if it needs to create or terminate a thread**



Kernel-Space Threads - Advantages

- *The disadvantages of User-space threads are the advantages of kernel-space threads*, for example:
- blocking system calls block individual threads, not processes
 - only the thread that wants to do IO is blocked, and *other threads in the process are free to run* and perform other activities
 - this allows *programs with significant IO to realize the largest increases in performance* through multi-threading (ie: there is the potential for increased CPU / IO overlap within processes)
- the kernel will schedule the threads individually
 - the *scheduler* is in the kernel, and *can pre-empt threads*
 - a single thread is *not able to monopolize the CPU*,
 - it is no longer necessary to call `thread_yield()`

enable maximum benefits of multi-core and multi-processor platforms

Kernel-Space Threads - Drawbacks

- thread management requires the program to *trap to the kernel*, which is *slower*
- thread support must be *built into the OS* which is costly and *technically challenging*
- there can be many, many, threads running concurrently on a modern system
 - requires a *lot of kernel memory*
 - kernel memory is expensive and (typically) pre-allocated
 - this increases the memory management burden within the kernel

Question 9

9. Which one of the following is an advantage of kernel space threads?
- a) superior performance for non-IO bound applications
 - b) fast thread switching
 - c) supports micro-kernel design concepts
 - d) blocking system calls do not affect other threads in the same process

Question 10

10. A program with multiple threads of execution that has hard real-time requirements () should **not** be developed on which of the following infrastructures:

- a) User-space threads
- b) Kernel-space threads
- c) Classic process model (no threads)
- d) it doesn't matter, any of the above approaches will work

What type of threads do different systems use?

- User-space threads
 - POSIX *Pthreads* (used on most traditional Unix platforms)
 - Solaris *threads*
- Kernel-space threads
 - Windows 95/98/NT/2000
 - Solaris
 - Linux
 - Mac OS-X 10.4 (Tiger)

Linux Threads

- Linux allows the programmer to manage the creation of threads, processes and/or both
 - note: *Linux refers to a thread as a task*
- to take full advantage of this you would have to completely understand the behaviour of your application and how it should operate
 - what kernel & OS info should be shared between processes
 - what kernel & OS info should be shared between tasks (threads)
- Linux adds a **new system call** to create threads and/or processes

```
pid = clone(function, stack_ptr, sharing_flags, args);
```


- *by setting sharing_flags appropriately you can create a new process, a new thread, or some hybrid entity*

Linux Threads - 2

Flag	Meaning when set	Meaning when cleared
CLONE_VM	Create a new thread	Create a new process
CLONE_FS	Share umask, root, and working dirs	Do not share them
CLONE_FILES	Share the file descriptors	Copy the file descriptors
CLONE_SIGHAND	Share the signal handler table	Copy the table
CLONE_PID	New thread gets old PID	New thread gets own PID
CLONE_PARENT	New thread has same parent as caller	New thread's parent is caller

- threads are supported in the kernel
- each thread will have a PID as well as a task ID
- there is a lot of flexibility with the above approach - but it will not port easily to other OS's
 - is this good? bad?

Solaris Threads



Java Threads

(from <http://www.cs.helsinki.fi/u/vihavain/k03/Java/Javathreads1.html>)

- threads are **implemented by the Java VM** (Virtual Machine)
 - their behaviour is *heavily influenced by the underlying operating system* and its characteristics
- the actual **scheduling policy** is system-dependent,
 - *determined together by the host OS and the VM implementation*
 - scheduling between multiple threads *may be preemptive* using time-slicing techniques, so that activities become (randomly) interleaved
 - *or it could be non-preemptive* where each thread must itself voluntarily give turn to others (yield)
- the behaviour of threads is **inherently nondeterministic**
 - ie, the *order of execution of threads is not repeatable* without explicit synchronization
 - Java *does not support hard real-time* systems programming (absolute timing requirements)

Java Threads - 2

- implementations can be either user-level threads package or package managed by the platform OS kernel
- "*green threads*" are simulated threads within the VM
- **green threads exist only at the user-level** and are not mapped to multiple kernel threads by the operating system
- "*native threads*" are the threads that are provided by the native OS
- native threads can realize the performance enhancement from parallelism (multiple CPUs)
- Java is naturally multi-threaded and because of this the **underlying OS implementation can make a substantial difference in the performance of your application**

Question 11

11. Assume you have a multi-threaded Java program running on Linux. The threads in your program will be managed by the JVM, which will map them to Pthreads in the kernel.

- a) TRUE
- b) FALSE

Justification:

Question 12

12. What does it mean if the `CLONE_VM` flag is set when you are making a `clone()` system call on Linux?
- a) a new processes should be created
 - b) copy the file descriptors
 - c) share the signal handler table
 - d) new thread gets its own files and resources
 - e) none of the above

The End