


# COMP 4735: Operating Systems

## Lesson 8.3: Semaphores



Rob Neilson

[rneilson@bcit.ca](mailto:rneilson@bcit.ca)



# Semaphores

# Cooperating Multiple Processes

- we use the `fork(proc, N, arg1, arg2, ..., argN)` command to create a process with the N arguments specified
- two cooperating processes can be defined and created as follows:

```
proc_0() {  
    while(TRUE) {  
        <compute section>;  
        <critical section>;  
    }  
}  
  
proc_1(int x, int y) {  
    while(TRUE) {  
        <compute section>;  
        <critical section>;  
    }  
}  
  
<shared global declarations>  
<initial processing>  
fork(proc_0, 0);  
fork(proc_1, 2, 6, 11);
```

- only one process at a time can be executing in its critical section
- each process also has a compute section, which may be executed concurrently with the other process

# What is a Semaphore?

- a semaphore,  $s$ , is a nonnegative integer variable that can only be changed or tested by these two indivisible functions:

```
P(s): [while(s == 0) {wait}; s = s - 1]
V(s): [s = s + 1]
```

down

P(s): accomplishes the “test and set” of a variable in a single indivisible instruction

up

V(s): signals any blocked processes to allow execution to continue

- these functions operate in a similar manner to the enter(lock) and exit(lock) functions described earlier, although they offer a more general solution

# To Solve a General Problem ...

```
Proc_0() {  
    while(TRUE) {  
        <compute section>;  
        P(s);  
        <critical section>;  
        V(s);  
    }  
}
```

```
proc_1() {  
    while(TRUE) {  
        <compute section>;  
        P(s);  
        <critical section>;  
        V(s);  
    }  
}
```

```
semaphore s = 1;  
fork(proc_0, 0);  
fork(proc_1, 0);
```

Insert **P(s)** immediately before the critical section.

- decrements the variable (s) so that next process has to wait

Insert **V(s)** immediately after the critical section.

- increments the variable (s) so that next process can continue

# Solving the Account Balance Problem

```
proc_0() {  
    . . .  
    P(s);  
    balance += amount;  
    V(s);  
    . . .  
}
```

```
proc_1() {  
    . . .  
    P(s);  
    balance -= amount;  
    V(s);  
    . . .  
}
```

```
semaphore s = 1;  
fork(proc_0, 0);  
fork(proc_1, 0);
```

*Remember:*

**P(s): [while(s == 0) {wait}; s = s - 1]**

**V(s): [s = s + 1]**




# Binary Semaphore

- the semaphore variable can be initialized to any non-negative value
- when a semaphore can only have values of 0 or 1, we call it a ***binary semaphore***
  - binary semaphore can be initialized to 0 or 1, depending on problem
    - most critical section problems initialize to 1
    - some synchronization problems initialize to 0 (see next example)
- the previous “critical section synchronization problem” was solved using a binary semaphore
- another general type of problem is the need to synchronize the actions of two threads, for example ....



# General Synchronization Problem

- consider the case where one process write a shared variable, and another reads it
- the process that is reading must wait for the data to be placed in the memory, or its read will obtain stale data



- in this problem there is ***no critical section***
- instead, the problem is to 'coordinate the actions' of the processes

# Formulating the General Sync Problem

we need this ...

```
proc_A() {  
  while(TRUE) {  
    <compute section A1>;  
    update(x);  
    <compute section A2>;  
    retrieve(y);  
  }  
}
```

```
fork(proc_A, 0);  
fork(proc_B, 0);
```

to happen before this ...

```
proc_B() {  
  while(TRUE) {  
    retrieve(x);  
    <compute section B1>;  
    update(y);  
    <compute section B2>;  
  }  
}
```

... and we need this ...

... to happen before this ...

# General Sync Problem (1)

```
proc_A() {  
    while(TRUE) {  
        <compute section A1>;  
        update(x);  
        <compute section A2>;  
        retrieve(y);  
    }  
}
```

```
proc_B() {  
    while(TRUE) {  
        retrieve(x);  
        <compute section B1>;  
        update(y);  
        <compute section B2>;  
    }  
}
```

```
semaphore s1 = 0; // proc_a will use this to tell proc_b  
                  // when it has written X
```

```
semaphore s2 = 0; // proc_b will use this to tell proc_a  
                  // when it has written Y
```

```
fork(proc_A, 0);  
fork(proc_B, 0);
```

- The solution will use two binary semaphores.
- The semaphores are initialized to 0 which means:
  - *wait until it is OK to proceed.*
- Binary semaphores that are initialized to 0 are sometimes called **blocking semaphores**

# General Sync Problem (2)

```
proc_A() {  
    while(TRUE) {  
        <compute section A1>;  
        update(x);  
        <compute section A2>;  
        retrieve(y);  
    }  
}
```

```
semaphore s1 = 0;  
semaphore s2 = 0;  
fork(proc_A, 0);  
fork(proc_B, 0);
```

```
proc_B() {  
    while(TRUE) {  
        P(s1); // wait for proc_a  
        retrieve(x);  
        <compute section B1>;  
        update(y);  
        <compute section B2>;  
    }  
}
```

Step 1:

- make proc\_b wait for proc\_a to write X
- s1 is initialized to 0, so proc b cannot proceed until explicitly told to
  - ie: it has to wait for proc\_a to do V(s1)

# General Sync Problem (3)

```
proc_A() {  
    while(TRUE) {  
        <compute section A1>;  
        update(x);  
        V(s1); // signal proc_b  
        <compute section A2>;  
        retrieve(y);  
    }  
}
```

```
proc_B() {  
    while(TRUE) {  
        P(s1); // wait for proc_a  
        retrieve(x);  
        <compute section B1>;  
        update(y);  
        <compute section B2>;  
    }  
}
```

```
semaphore s1 = 0;  
semaphore s2 = 0;  
fork(proc_A, 0);  
fork(proc_B, 0);
```

Step 2:

- make proc\_a signal proc\_b after it writes X
- V(s1) increments the semaphore to 1
- this allows proc\_b to proceed the next time it gets the CPU

# General Sync Problem (4)

```
proc_A() {  
    while(TRUE) {  
        <compute section A1>;  
        update(x);  
        V(s1); // signal proc_b  
        <compute section A2>;  
        P(s2); // wait for proc_b  
        retrieve(y);  
    }  
}
```

```
proc_B() {  
    while(TRUE) {  
        P(s1); // wait for proc_a  
        retrieve(x);  
        <compute section B1>;  
        update(y);  
        <compute section B2>;  
    }  
}
```

```
semaphore s1 = 0;  
semaphore s2 = 0;  
fork(proc_A, 0);  
fork(proc_B, 0);
```

Step 3:

- make proc\_a wait for proc\_b to write Y
- note that we are using the other semaphore
  - ie: P(s2)



# General Sync Problem (5)

```
proc_A() {  
    while(TRUE) {  
        <compute section A1>;  
        update(x);  
        V(s1); // signal proc_b  
        <compute section A2>;  
        P(s2); // wait for proc_b  
        retrieve(y);  
    }  
}
```

```
proc_B() {  
    while(TRUE) {  
        P(s1); // wait for proc_a  
        retrieve(x);  
        <compute section B1>;  
        update(y);  
        V(s2); // signal proc_a  
        <compute section B2>;  
    }  
}
```

```
semaphore s1 = 0;  
semaphore s2 = 0;  
fork(proc_A, 0);  
fork(proc_B, 0);
```

Step 4:

- make proc\_b signal proc\_a after it writes Y
- V(s2) increments the semaphore to 1
- this allows proc\_a to proceed the next time it gets the CPU

# Final Solution to General Problem

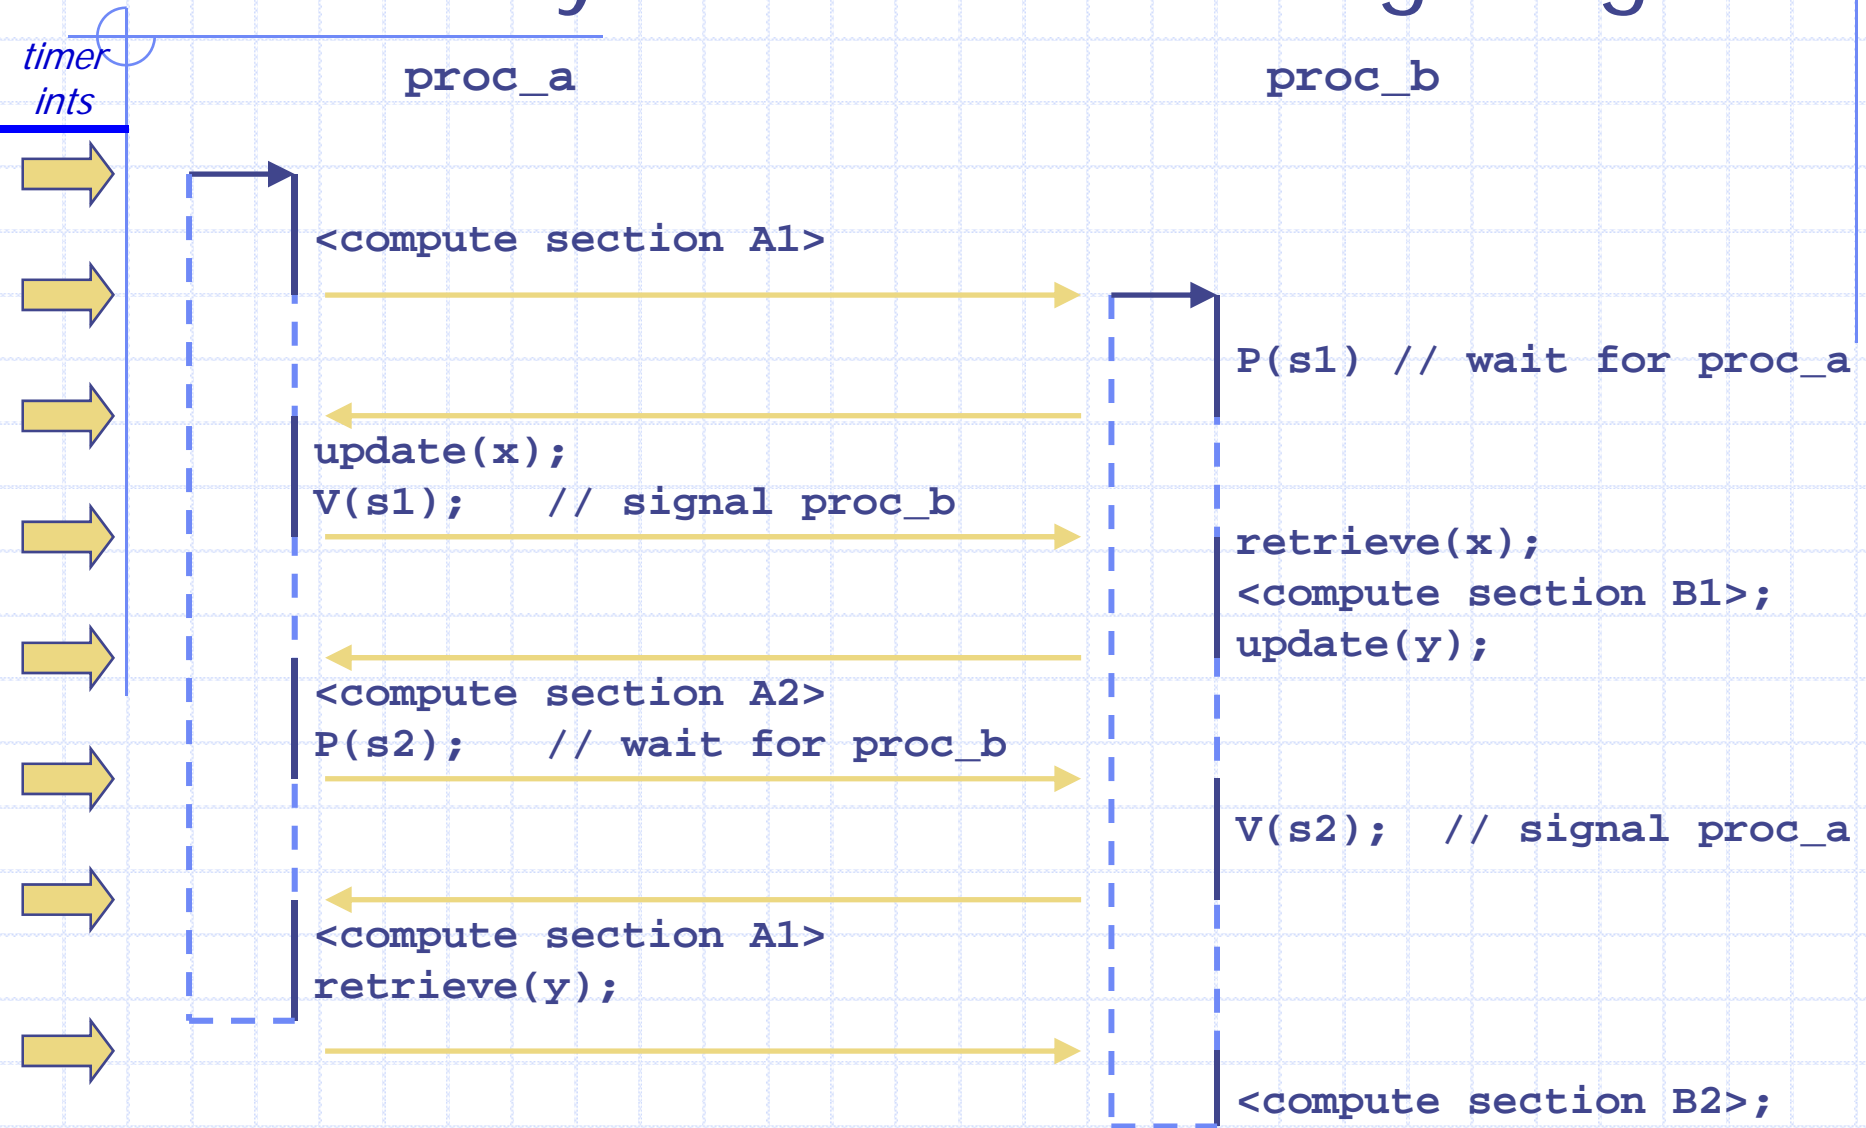
```
proc_A() {  
    while(TRUE) {  
        <compute section A1>;  
        update(x);  
        V(s1); // signal proc_b  
        <compute section A2>;  
        P(s2); // wait for proc_b  
        retrieve(y);  
    }  
}
```

```
proc_B() {  
    while(TRUE) {  
        P(s1); // wait for proc_a  
        retrieve(x);  
        <compute section B1>;  
        update(y);  
        V(s2); // signal proc_a  
        <compute section B2>;  
    }  
}
```

```
semaphore s1 = 0;  
semaphore s2 = 0;  
fork(proc_A, 0);  
fork(proc_B, 0);
```



# General Sync Problem: Timing Diag



# Counting Semaphore

So far we have considered two types of semaphores:

1. **binary semaphore**:
  - used to control access to critical sections
2. **blocking semaphores**:
  - used to make threads wait for events before proceeding

There is a more general type of semaphore, called a **counting semaphore**, which can take on values other than 0 or 1.

- counting semaphores are useful for situations where you want to limit the number of threads that are in some state, but the limit is greater than 1
  - for example, to access or use a resource

# Counting Semaphores: Resource Control

## Problem:

- you have a finite amount of some resource
- many threads can acquire and use a resource
- if more threads request resources than there are resources, we want threads to block until a resource becomes available

## Solution:

- create a **counting semaphore**, shared among all threads
- **initialize** it to the **number of resources**
- **when a thread wants a resource it executes  $P(s)$**
- if  $s$  is 0 (no resources available) the thread blocks
- if  $s > 0$ ,  $s$  is decremented and one of the resources is assigned to the thread
- **when a thread is done with the resource it executes  $V(s)$  to increment the semaphore (put the resource back in the pool)**

# Resource Control Solution

```
proc() {  
    while(TRUE) {  
        <compute section>;  
        P(res); // block if no resources available  
        <use the resource>  
        V(res); // free a resource  
    }  
}  
  
semaphore res = 5; // there are 5 resources to share  
for (i=1 to 100)  
    fork(proc, 0); // there are 100 processes
```

# Classic Semaphore Problems

- there are a number of classic semaphore problems that are worth studying as they provide insight into how semaphores can be used
- these problems include
  - bounded-buffer (aka producer-consumer) problem
  - readers-writers problem
  - dining philosophers problem
- we will consider some of these problems now ...

# Bounded Buffer Problem Statement

There are two threads that share data using  $N$  buffers.

- the buffers exist in and are accessed via shared memory

## 1. producer thread operates as follows ...

- get an empty buffer
- write some information into the buffer
- put the full buffer somewhere that the consumer process can get it


Note: the producer must block if there are no empty buffers to fill

## 2. consumer thread operates as follows ...

- get a full buffer
- extract (and process) the information in it
- put the empty buffer somewhere that the producer process can get it

Note: the consumer must block if there are no full buffers to process

# Bounded Buffer Scenario





# Solving the Bounded Buffer Problem

To solve the problem, we need to consider all operations that need to be synchronized.

– ie: where will we use semaphores?

1. producer needs to **block** if there are **no empty buffers** to fill
  - we will use a **counting semaphore**, initialized to the **number of empty buffers**  $N$
  - each time the producer takes a buffer, the semaphore will be decremented
  - producer blocks if semaphore is 0 when it tries to take a buffer
2. consumer needs to **block** if there are **no full buffers** to process
  - we will use a **counting semaphore**, initialized to the **number of full buffers** 0
  - each time the consumer gets a full buffer, the semaphore will be decremented
  - consumer blocks if semaphore is 0 when it tries to get a buffer to process
3. both threads need to enter **critical sections** when they update the contents of the buffer pools
  - we need to prevent race conditions, as the buffer pools are shared variables
  - we will use a **binary semaphore** to ensure that only **one thread** updates the contents of the buffer pool **at a time**



# Bounded Buffer Problem - Producer

```
producer() {
    buf_type *next, *here;
    while(TRUE) {
        produce_item(next);
        P(empty);
        P(s);
        here = obtain(empty);
        V(s);
        copy_buffer(next, here);
        P(s);
        release(here, fullPool);
        V(s);
        V(full);
    }
}

semaphore s = 1;
semaphore full = 0;
semaphore empty = N;
buf_type buffer[N];
fork(producer, 0);
fork(consumer, 0);

// next holds the data to write
// create some data for consumer
// get empty buff; block if none
// enter critical section
// remove buff from empty pool
// exit the critical section
// copy 'next' into the new buffer
// enter critical section
// add buffer to full pool
// exit critical section
// signal consumer that there is...
// ... data to process

// binary semaphore for mutual excl
// counting semaphore for full pool
// counting semaphore for empty pool
```

# Bounded Buffer Problem - Consumer

```
consumer() {
    buf_type *next, *here;
    while(TRUE) {
        P(full);                // get full buff; block if none
        P(s);                   // enter critical section
        here = obtain(full);    // get buff from full pool
        V(s);                   // exit the critical section
        copy_buffer(here, next); // empty the buffer (to local buff)
        P(s);                   // enter critical section
        release(here, emptyPool); // put buffer back in empty pool
        V(s);                   // exit critical section
        V(empty);               // signal producer (buffs available)
        consumeItem(next);      // use the data for whatever ...
    }                           // ... purpose it was intended
}

semaphore s = 1;
semaphore full = 0;
semaphore empty = N;
buf_type buffer[N];
fork(producer, 0);
fork(consumer, 0);
```

# Bounded Buffer Problem (final soln)

```
producer() {  
    buf_type *next, *here;  
    while(TRUE) {  
        produce_item(next);  
        P(empty);  
        P(s);  
        here = obtain(empty);  
        V(s);  
        copy_buffer(next, here);  
        P(s);  
        release(here, fullPool);  
        V(s);  
        V(full);  
    }  
}
```

```
semaphore s = 1;  
semaphore full = 0;  
semaphore empty = N;  
buf_type buffer[N];  
fork(producer, 0);  
fork(consumer, 0);
```

```
consumer() {  
    buf_type *next, *here;  
    while(TRUE) {  
        P(full);  
        P(s);  
        here = obtain(full);  
        V(s);  
        copy_buffer(here, next);  
        P(s);  
        release(here, emptyPool);  
        V(s);  
        V(empty);  
        consumeItem(next);  
    }  
}
```



green lines are not protected; just general processing



blue lines are counting/synchronising buffer use



red lines are protected critical sections

# The End