


COMP 4735: Operating Systems Concepts

Lesson 12: Page Replacement Algorithms



Rob Neilson

rneilson@bcit.ca

Administrative stuff ...

- Reading: Virtual Memory is covered in chapters 3.3, 3.4
- Next quiz: Monday March 30th
 - covers chapters 3.1, 3.2, 3.3 plus all stuff from lecture
 - note: chapter 3.4 will be covered in a later quiz



Paging Algorithms

Paging Algorithms and Policies ...

Any paging algorithm must define:

- fetch policy
 - decides when a page should be loaded into memory
- replacement policy
 - decides which page should be moved to disk when all pages are full
- placement policy
 - decides where the page should be loaded in primary memory

There are two types of paging algorithms:

- static paging algorithms
 - each process uses a fixed number of frames during its lifetime
- dynamic paging algorithms
 - the number of frames allocated to a process changes during its lifetime

Static Paging Algorithms

- number of page frames does not change throughout lifetime of the process
- this makes the *placement policy* rather simple ...
 - assume that all of a processes page frames are full (this is true at all times except when process first starts)
 - assume a new page is needed
 - the replacement policy is invoked to select a victim, and the selected page is moved to disk
 - the new page is loaded into the frame that was just vacated
- this is always the pattern for static paging algorithms
- in dynamic paging algorithms we have other choices, such as to grow the number of frames allocated to the process

Demand Paging

- “Demand Paging” is a *fetch policy*
 - this is the typical policy that says to load a page only when it has been referenced
 - ie: this is what we have been talking about in previous slides/lectures
- What are the alternatives?
 - we could try to predict when a page will be needed, and load it ahead of time
 - in this case we are essentially performing a “prefetch”
- Static Paging Algorithms all use the **demand paging** fetch policy
- Static Paging Algorithms all use the simple **fetch policy** described on the previous page

*this means that we only need to consider **replacement policies** in our study of static paging algorithms*

Static Page Frame Allocation with Demand Paging

The following Static Page Replacement algos will be considered:

- The Optimal Algorithm
- First In First Out (FIFO)
- Least Recently Used (LRU)
- Not Frequently Used (NFU)

Background

- In looking at these algorithms, we will consider a continuous stream of page requests R . For example:

$R = 0, 1, 2, 3, 0, 1, 2, 3, 0, 1, 2, 3, 4, 5, 6$

- The above example tells us that page 0 is requested first, then page 1, and so on.
- We also consider the state S of the system at any point in time.
 - the state is just a snapshot of the pages that are currently loaded. For example:

$S = \{ 3, 6, 9, 2 \}$

- this tells us that there are 4 available frames, and in the current state:
 - frame 0 holds page 3
 - frame 1 holds page 6
 - frame 2 holds page 9
 - frame 3 holds page 2

Diagramming Page Replacement - 1

Assume we are given:

- a reference stream R (eg: $R = 3, 2, 0, 7, 0, 3, 4, 1, 4, 5$)
- an initial state S (eg: $S = \{ \}$)
- a replacement policy (eg: FIFO – first page in is next one out)

We are now able to illustrate how the page replacement policy works, for example:

Request			3	2	0	7	0	3	4	1	4	5
Frame	0	-										
	1	-										
	2	-										
	3	-										
Page faults												

Diagramming Page Replacement - 2

Assume we are given:

- a reference stream R (eg: $R = 3, 2, 0, 7, 0, 3, 4, 1, 4, 5$)
- an initial state S (eg: $S = \{ \}$)
- a replacement policy (eg: FIFO – first page in is next one out)

We are now able to illustrate how the page replacement policy works, for example:

Request			3	2	0	7	0	3	4	1	4	5
Frame	0	-	<u>3</u>									
	1	-										
	2	-										
	3	-										
Page faults			*									

Diagramming Page Replacement - 3

Assume we are given:

- a reference stream R (eg: R = 3, 2, 0, 7, 0, 3, 4, 1, 4, 5)
- an initial state S (eg: S = { })
- a replacement policy (eg: FIFO – first page in is next one out)

We are now able to illustrate how the page replacement policy works, for example:

Request			3	2	0	7	0	3	4	1	4	5
Frame	0	-	<u>3</u>	3								
	1	-		<u>2</u>								
	2	-										
	3	-										
Page faults			*	*								

Diagramming Page Replacement - 4

Assume we are given:

- a reference stream R (eg: $R = 3, 2, 0, 7, 0, 3, 4, 1, 4, 5$)
- an initial state S (eg: $S = \{ \}$)
- a replacement policy (eg: FIFO – first page in is next one out)

We are now able to illustrate how the page replacement policy works, for example:

Request			3	2	0	7	0	3	4	1	4	5
Frame	0	-	<u>3</u>	3	3							
	1	-		<u>2</u>	2							
	2	-			<u>0</u>							
	3	-										
Page faults			*	*	*							

Diagramming Page Replacement - 5

Assume we are given:

- a reference stream R (eg: R = 3, 2, 0, 7, 0, 3, 4, 1, 4, 5)
- an initial state S (eg: S = { })
- a replacement policy (eg: FIFO – first page in is next one out)

We are now able to illustrate how the page replacement policy works, for example:

Request			3	2	0	7	0	3	4	1	4	5
Frame	0	-	<u>3</u>	3	3	3						
	1	-		<u>2</u>	2	2						
	2	-			<u>0</u>	0						
	3	-				<u>7</u>						
Page faults			*	*	*	*						

Diagramming Page Replacement - 6

Assume we are given:

- a reference stream R (eg: R = 3, 2, 0, 7, 0, 3, 4, 1, 4, 5)
- an initial state S (eg: S = { })
- a replacement policy (eg: FIFO – first page in is next one out)

We are now able to illustrate how the page replacement policy works, for example:

Request			3	2	0	7	0	3	4	1	4	5
Frame	0	-	<u>3</u>	3	3	3	3					
	1	-		<u>2</u>	2	2	2					
	2	-			<u>0</u>	0	0					
	3	-				<u>7</u>	7					
Page faults			*	*	*	*						

Diagramming Page Replacement - 7

Assume we are given:

- a reference stream R (eg: R = 3, 2, 0, 7, 0, 3, 4, 1, 4, 5)
- an initial state S (eg: S = { })
- a replacement policy (eg: FIFO – first page in is next one out)

We are now able to illustrate how the page replacement policy works, for example:

Request			3	2	0	7	0	3	4	1	4	5
Frame	0	-	<u>3</u>	3	3	3	3	3				
	1	-		<u>2</u>	2	2	2	2				
	2	-			<u>0</u>	0	0	0				
	3	-				<u>7</u>	7	7				
Page faults			*	*	*	*						

Diagramming Page Replacement - 8

Assume we are given:

- a reference stream R (eg: R = 3, 2, 0, 7, 0, 3, 4, 1, 4, 5)
- an initial state S (eg: S = { })
- a replacement policy (eg: FIFO – first page in is next one out)

We are now able to illustrate how the page replacement policy works, for example:

Request			3	2	0	7	0	3	4	1	4	5
Frame	0	-	<u>3</u>	3	3	3	3	3	<u>4</u>			
	1	-		<u>2</u>	2	2	2	2	2			
	2	-			<u>0</u>	0	0	0	0			
	3	-				<u>7</u>	7	7	7			
Page faults			*	*	*	*			*			

Diagramming Page Replacement - 9

Assume we are given:

- a reference stream R (eg: R = 3, 2, 0, 7, 0, 3, 4, 1, 4, 5)
- an initial state S (eg: S = { })
- a replacement policy (eg: FIFO – first page in is next one out)

We are now able to illustrate how the page replacement policy works, for example:

Request			3	2	0	7	0	3	4	1	4	5
Frame	0	-	<u>3</u>	3	3	3	3	3	<u>4</u>	4		
	1	-		<u>2</u>	2	2	2	2	2	<u>1</u>		
	2	-			<u>0</u>	0	0	0	0	0		
	3	-				<u>7</u>	7	7	7	7		
Page faults			*	*	*	*			*	*		

Diagramming Page Replacement - 10

Assume we are given:

- a reference stream R (eg: R = 3, 2, 0, 7, 0, 3, 4, 1, 4, 5)
- an initial state S (eg: S = { })
- a replacement policy (eg: FIFO – first page in is next one out)

We are now able to illustrate how the page replacement policy works, for example:

Request			3	2	0	7	0	3	4	1	4	5
Frame	0	-	<u>3</u>	3	3	3	3	3	<u>4</u>	4	4	
	1	-		<u>2</u>	2	2	2	2	2	<u>1</u>	1	
	2	-			<u>0</u>	0	0	0	0	0	0	
	3	-				<u>7</u>	7	7	7	7	7	
Page faults			*	*	*	*			*	*		

Diagramming Page Replacement - 11

Assume we are given:

- a reference stream R (eg: R = 3, 2, 0, 7, 0, 3, 4, 1, 4, 5)
- an initial state S (eg: S = { })
- a replacement policy (eg: FIFO – first page in is next one out)

We are now able to illustrate how the page replacement policy works, for example:

Request			3	2	0	7	0	3	4	1	4	5
Frame	0	-	<u>3</u>	3	3	3	3	3	<u>4</u>	4	4	4
	1	-		<u>2</u>	2	2	2	2	2	<u>1</u>	1	1
	2	-			<u>0</u>	0	0	0	0	0	0	<u>5</u>
	3	-				<u>7</u>	7	7	7	7	7	7
Page faults			*	*	*	*			*	*		*

The optimal page replacement algorithm - 1

Idea:

- Select the page that will not be needed for the longest time
 - We do this by looking at the “forward distance” to the page
 - this is the distance from the current point in the reference stream to the next place in the stream where the same page is referenced
 - forward distance always > 0
 - forward distance $= \infty$ if the page is never referenced again
 - In the case of ties, choose arbitrarily
-
- Consider a reference stream:

$R = 3, 2, 0, 5, 4, 3, 1, 5, 2, 0$

4 frames, all empty to start ...

The optimal page replacement algorithm - 2

Idea:

- Select the page that will not be needed for the longest time
- We do this by looking at the “forward distance” to the page
 - this is the distance from the current point in the reference stream to the next place in the stream where the same page is referenced
 - forward distance always > 0
 - forward distance $== \infty$ if the page is never referenced again
- In the case of ties, choose arbitrarily

Request			3	2	0	5	4	3	1	5	2	0
Frame	0	-	<u>3</u>	3	3	3						
	1	-		<u>2</u>	2	2						
	2	-			<u>0</u>	0						
	3	-				<u>5</u>						
Page faults			*	*	*	*						

The optimal page replacement algorithm - 3

Idea:

- Select the page that will not be needed for the longest time
- We do this by looking at the “forward distance” to the page
 - this is the distance from the current point in the reference stream to the next place in the stream where the same page is referenced
 - forward distance always > 0
 - forward distance $= \infty$ if the page is never referenced again
- In the case of ties, choose arbitrarily

Request			3	2	0	5	4	3	1	5	2	0
Frame	0	–	<u>3</u>	3	3	3	<i>forward distance = 1</i>					
	1	–		<u>2</u>	2	2	<i>forward distance = 4</i>					
	2	–			<u>0</u>	0	<i>forward distance = 5</i>					
	3	–				<u>5</u>	<i>forward distance = 3</i>					
Page faults			*	*	*	*						

current point

The optimal page replacement algorithm - 4

Idea:

- Select the page that will not be needed for the longest time
- We do this by looking at the “forward distance” to the page
 - this is the distance from the current point in the reference stream to the next place in the stream where the same page is referenced
 - forward distance always > 0
 - forward distance $== \infty$ if the page is never referenced again
- In the case of ties, choose arbitrarily

Request			3	2	0	5	4	3	1	5	2	0
Frame	0	–	<u>3</u>	3	3	3	3					
	1	–		<u>2</u>	2	2	2					
	2	–			<u>0</u>	0	<u>4</u>					
	3	–				<u>5</u>	5					
Page faults			*	*	*	*	*					

The optimal page replacement algorithm - 5

Idea:

- Select the page that will not be needed for the longest time
- We do this by looking at the “forward distance” to the page
 - this is the distance from the current point in the reference stream to the next place in the stream where the same page is referenced
 - forward distance always > 0
 - forward distance $== \infty$ if the page is never referenced again
- In the case of ties, choose arbitrarily

Request			3	2	0	5	4	3	1	5	2	0
Frame	0	-	<u>3</u>	3	3	3	3	3				
	1	-		<u>2</u>	2	2	2	2				
	2	-			<u>0</u>	0	<u>4</u>	4				
	3	-				<u>5</u>	5	5				
Page faults			*	*	*	*	*					

The optimal page replacement algorithm - 6

Idea:

- Select the page that will not be needed for the longest time
- We do this by looking at the “forward distance” to the page
 - this is the distance from the current point in the reference stream to the next place in the stream where the same page is referenced
 - forward distance always > 0
 - forward distance $== \infty$ if the page is never referenced again
- In the case of ties, choose arbitrarily

Request			3	2	0	5	4	3	1	5	2	0
Frame	0	–	<u>3</u>	3	3	3	3	3	<u>1</u>			
	1	–		<u>2</u>	2	2	2	2	2			
	2	–			<u>0</u>	0	<u>4</u>	4	4			
	3	–				<u>5</u>	5	5	5			
Page faults			*	*	*	*	*		*			

The optimal page replacement algorithm - 7

Idea:

- Select the page that will not be needed for the longest time
- We do this by looking at the “forward distance” to the page
 - this is the distance from the current point in the reference stream to the next place in the stream where the same page is referenced
 - forward distance always > 0
 - forward distance $= \infty$ if the page is never referenced again
- In the case of ties, choose arbitrarily

Request			3	2	0	5	4	3	1	5	2	0
Frame	0	–	<u>3</u>	3	3	3	3	3	<u>1</u>	1		
	1	–		<u>2</u>	2	2	2	2	2	2		
	2	–			<u>0</u>	0	<u>4</u>	4	4	4		
	3	–				<u>5</u>	5	5	5	5		
Page faults			*	*	*	*	*		*			

The optimal page replacement algorithm - 8

Idea:

- Select the page that will not be needed for the longest time
- We do this by looking at the “forward distance” to the page
 - this is the distance from the current point in the reference stream to the next place in the stream where the same page is referenced
 - forward distance always > 0
 - forward distance $= \infty$ if the page is never referenced again
- In the case of ties, choose arbitrarily

Request			3	2	0	5	4	3	1	5	2	0
Frame	0	–	<u>3</u>	3	3	3	3	3	<u>1</u>	1	1	
	1	–		<u>2</u>	2	2	2	2	2	2	2	
	2	–			<u>0</u>	0	<u>4</u>	4	4	4	4	
	3	–				<u>5</u>	5	5	5	5	5	
Page faults			*	*	*	*	*		*			

The optimal page replacement algorithm - 9

Idea:

- Select the page that will not be needed for the longest time
- We do this by looking at the “forward distance” to the page
 - this is the distance from the current point in the reference stream to the next place in the stream where the same page is referenced
 - forward distance always > 0
 - forward distance $= \infty$ if the page is never referenced again
- In the case of ties, choose arbitrarily

Request			3	2	0	5	4	3	1	5	2	0
Frame	0	-	<u>3</u>	3	3	3	3	3	<u>1</u>	1	1	?
	1	-		<u>2</u>	2	2	2	2	2	2	2	?
	2	-			<u>0</u>	0	<u>4</u>	4	4	4	4	?
	3	-				<u>5</u>	5	5	5	5	5	?
Page faults			*	*	*	*	*		*			*

Optimal page replacement – example 2

Input:

- $R = 0, 1, 2, 3, 0, 1, 2, 3, 0, 1, 2, 3, 4, 5, 6, 7$
- 3 frames, all start empty

Try and work through this one yourself.

Request		0	1	2	3	0	1	2	3	0	1	2	3	4	5	6	7
Frame	0																
	1																
	2																
Page faults																	

Optimal page replacement – example 2

Input:

- R = 0, 1, 2, 3, 0, 1, 2, 3, 0, 1, 2, 3, 4, 5, 6, 7
- 3 frames, all start empty

Try and work through this one yourself.

Request			0	1	2	3	0	1	2	3	0	1	2	3	4	5	6	7
Frame	0	-	<u>0</u>	0	0	0	0	0	0	0	0	<u>1</u>	1	1	<u>4</u>	4	4	<u>7</u>
	1	-		<u>1</u>	1	1	1	<u>1</u>	<u>2</u>	2	2	2	<u>2</u>	2	2	<u>5</u>	5	5
	2	-			<u>2</u>	<u>3</u>	3	3	3	3	3	3	3	3	3	3	<u>6</u>	6
Page faults			*	*	*	*			*			*			*	*	*	*

The optimal page replacement algorithm

Idea:

- Select the page that will not be needed for the longest time

Problem:

- Can't know the future of a program
- Can't know when a given page will be needed next
- The optimal algorithm is unrealizable

However:

- We can use it only as a control case for simulation studies
 - Run the program once
 - Generate a log of all memory references
 - Use the log to simulate various page replacement algorithms
 - Can compare other algorithms to this “optimal” algorithm
 - for example, we know that the last example included 10 page faults under the optimal algorithm

In-class Exercise: Optimal Algorithm

- Assume that a memory system uses the optimal algorithm.
- The system has 4 page frames, and they are all empty.
- Given the following reference stream, calculate the number of page faults that occur.

$R = 1\ 3\ 9\ 2\ 4\ 3\ 9\ 1\ 8\ 5\ 4\ 5\ 4\ 3\ 9\ 1\ 8\ 2\ 3\ 9\ 8\ 1\ 2\ 3$

- I am not going to work through the answer, but I will tell you what it is.
 - the answer is: 11 page faults

FIFO page replacement algorithm

- Always replace the oldest page ...
 - *"Replace the page that has been in memory for the longest time."*
 - We used this technique in the earlier explanation of reference streams
- It is *not* a good algorithm, as it makes no assumptions about when a page was used or is about to be used
- Consider the reference streams from the previous example ...

Request			0	1	2	3	0	1	2	3	0	1	2	3	4	5	6	7
Frame	0	-	<u>0</u>	0	0	<u>3</u>	3	3	<u>2</u>	2	2	<u>1</u>	1	1	<u>4</u>	4	4	<u>7</u>
	1	-		<u>1</u>	1	1	<u>0</u>	0	0	<u>3</u>	3	3	<u>2</u>	2	2	<u>5</u>	5	5
	2	-			<u>2</u>	2	2	<u>1</u>	1	1	<u>0</u>	0	0	<u>3</u>	3	3	<u>6</u>	6
Page faults			*	*	*	*	*	*	*	*	*	*	*	*	*	*	*	*

FIFO page replacement algorithm

- This time we had 16 page faults (compared to 10 with optimal algorithm)
- The problem is that the replaced page may be needed again soon

Request			0	1	2	3	0	1	2	3	0	1	2	3	4	5	6	7
Frame	0	-	<u>0</u>	0	0	<u>3</u>	3	3	<u>2</u>	2	2	<u>1</u>	1	1	<u>4</u>	4	4	<u>7</u>
	1	-		<u>1</u>	1	1	<u>0</u>	0	0	<u>3</u>	3	3	<u>2</u>	2	2	<u>5</u>	5	5
	2	-			<u>2</u>	2	2	<u>1</u>	1	1	<u>0</u>	0	0	<u>3</u>	3	3	<u>6</u>	6
Page faults			*	*	*	*	*	*	*	*	*	*	*	*	*	*	*	*

FIFO page replacement algorithm

- And with the other reference stream used earlier ...
- This stream gives 9 pages faults, while optimal is 6.

Request			3	2	0	5	4	3	1	5	2	0
Frame	0	-	<u>3</u>	3	3	3	<u>4</u>	4	4	4	4	<u>0</u>
	1	-		<u>2</u>	2	2	2	<u>3</u>	3	3	3	3
	2	-			<u>0</u>	0	0	0	<u>1</u>	1	1	1
	3	-				<u>5</u>	5	5	5	5	<u>2</u>	2
Page faults			*	*	*	*	*	*	*		*	*

In-class Exercise: FIFO Algorithm

- Assume that a memory system uses the FIFO algorithm.
- The system has 4 page frames, and they are all empty.
- Given the following reference stream, calculate the number of page faults that occur.

R = 1 3 9 2 4 3 9 1 8 5 4 5 4 3 9 1 8 2 3 9 8 1 2 3

– the answer is: 16 page faults

Least recently used algorithm (LRU)

- Keep track of when a page is used
- Replace the page that has been used least recently
 - ie: replace the page that hasn't been referenced in the longest time
 - essentially, we use the "largest backward distance" to select pages to replace
- This algorithm should do better, as it is exploiting "locality of reference"
... let's see how it does ...

Request			0	1	2	3	0	1	2	3	0	1	2	3	4	5	6	7
Frame	0	-																
	1	-																
	2	-																
Page faults																		

Least recently used algorithm (LRU)

- 16 page faults again!
- This sucks ... but in this case we simply have a 'degenerative case'
- In other cases this algorithm has proven to do better than FIFO

Let's try the other reference stream ...

Request			0	1	2	3	0	1	2	3	0	1	2	3	4	5	6	7
Frame	0	-	<u>0</u>	0	0	<u>3</u>	3	3	<u>2</u>	2	2	<u>1</u>	1	1	<u>4</u>	4	4	<u>7</u>
	1	-		<u>1</u>	1	1	<u>0</u>	0	0	<u>3</u>	3	3	<u>2</u>	2	2	<u>5</u>	5	5
	2	-			<u>2</u>	2	2	<u>1</u>	1	1	<u>0</u>	0	0	<u>3</u>	3	3	<u>6</u>	6
Page faults			*	*	*	*	*	*	*	*	*	*	*	*	*	*	*	*

LRU page replacement algorithm

- Again we get 9 pages faults, while optimal is 6 – same as FIFO!
- This sucks as well.
- But it is a better algorithm than FIFO. Have a browse through 3.4.6 in your text if you want some more information...

Request			3	2	0	5	4	3	1	5	2	0
Frame	0	-	<u>3</u>	3	3	3	<u>4</u>	4	4	4	<u>2</u>	2
	1	-		<u>2</u>	2	2	2	<u>3</u>	3	3	3	<u>0</u>
	2	-			<u>0</u>	0	0	0	<u>1</u>	1	1	1
	3	-				<u>5</u>	5	5	5	5	5	5
Page faults			*	*	*	*	*	*	*		*	*

In-class Exercise: LRU Algorithm

- Assume that a memory system uses the LRU algorithm.
- The system has 4 page frames, and they are all empty.
- Given the following reference stream, calculate the number of page faults that occur.

R = 1 3 9 2 4 3 9 1 8 5 4 5 4 3 9 1 8 2 3 9 8 1 2 3

- the answer is: 19 page faults

Not *frequently* used algorithm (NFU)

- Associate a counter with each page
 - Every time the page is referenced we increment the counter
 - The counter approximates how often the page is used
 - For replacement, choose the page with lowest counter
 - in the case of ties, select an arbitrary page to replace

*now we have a tie – all have been used once
so we will choose randomly ...*

Request			0	1	2	3	0	1	2	3	0	1	2	3	4	5	6	7
Frame	0	-	<u>0</u>	0	0													
	1	-		<u>1</u>	1													
	2	-			<u>2</u>													
Page faults			*	*	*													

Not frequently used algorithm (NFU)

- So we get 12 faults in total ... better than FIFO ... but only because we made a good arbitrary pick!
- It is interesting to notice what happens at the end, when a sequence of previously unused frames are loaded ...
... in this case the least frequently used frame is the last one loaded

Request			0	1	2	3	0	1	2	3	0	1	2	3	4	5	6	7
Frame	0	-	<u>0</u>	0	0	0	0	0	0	0	0	0	0	<u>3</u>	3	3	3	3
	1	-		<u>1</u>	1	1	1	1	1	<u>3</u>	3	<u>1</u>	1	1	1	1	1	1
	2	-			<u>2</u>	<u>3</u>	3	3	<u>2</u>	2	2	2	2	2	<u>4</u>	<u>5</u>	<u>6</u>	<u>7</u>
Page faults			*	*	*	*			*	*		*		*	*	*	*	*

Not frequently used algorithm (NFU)

- With our other reference stream we get 8.
- Better than LRU, but not as good as optimal (6).

Request			3	2	0	5	4	3	1	5	2	0
Frame	0	-	<u>3</u>	3	3	3	3	3	3	3	3	3
	1	-		<u>2</u>	2	2	2	2	2	2	2	2
	2	-			<u>0</u>	0	0	0	<u>1</u>	1	1	<u>0</u>
	3	-				<u>5</u>	<u>4</u>	4	4	<u>5</u>	5	5
Page faults			*	*	*	*	*		*	*		*

In-class Exercise: NFU Algorithm

- Assume that a memory system uses the NFU algorithm.
- The system has 4 page frames, and they are all empty.
- Given the following reference stream, calculate the number of page faults that occur.

R = 1 3 9 2 4 3 9 1 8 5 4 5 4 3 9 1 8 2 3 9 8 1 2 3

- the answer is: 14 page faults (depending on how you break ties)

Not frequently used algorithm (NFU)

Problem with NFU:

- Some page may be heavily used
 - ---> Its counter is large
- The program's behavior changes
 - Now, this page is not used ever again (or only rarely), but it stays in memory
- This algorithm never forgets!
 - *This page will never be chosen for replacement!*

Modified NFU with aging

- Associate a counter with each page like before
- Also keep a “time window” – so that we only consider the frequency of use within some previous period of time
- One way to do this is:
 - On every clock tick, the OS looks at each page.
 - Shift the counter right 1 bit (divide its value by 2)
 - If the **Page Reference Bit** is set...
 - Set the **most-significant bit**
 - Clear the Referenced Bit

100000 = 32

010000 = 16

001000 = 8

000100 = 4

100010 = 34

110001 = 49

111000 = 56

011100 = 28



Dynamic Paging Algorithms

Paging Algorithms and Policies (review)

Any paging algorithm must define:

- fetch policy
 - decides when a page should be loaded into memory
- replacement policy
 - decides which page should be moved to disk when all pages are full
- placement policy
 - decides where the page should be loaded in primary memory

There are two types of paging algorithms:

- static paging algorithms
 - each process uses a fixed number of frames during its lifetime
- dynamic paging algorithms
 - the number of frames allocated to a process changes during its lifetime

Static Paging Algorithms (review)

- number of page frames does not change throughout lifetime of the process
 - ie: each process is allocated a number of page frames, even if it doesn't need them, or if it leads to thrashing
- “thrashing” is the situation where there are not enough page frames allocated to the process, and it continually keeps swapping pages
 - this can occur when a loop crosses a page boundary, say from page n to $n+1$, but there is only one frame available
 - depending on replacement policy, n and $n+1$ could continue to displace each other as the program loops
 - this is very slow, as we have to write pages to disk on each iteration of the loop
- in this case it makes more sense to allocate an extra frame to the process for the times that it needs the extra page

Dynamic Paging Algorithms

- dynamic paging algorithms adjust the number of frames used by a process according to its need for these frames
- these algorithms specify both the placement and the replacement policies, ie:
 - they specify where the page can be loaded
 - they specify which page is to be replaced if there is a choice
- the motivation for dynamic paging algorithms is the concept of “locality of reference”
 - this is the notion that programs tend stay in specific memory locales within the process
 - for example, the main body of a program name be a loop that calls a few methods from within a loop.
 - in this case, the pages that contain the loop code and the most commonly accessed methods will be accessed much more frequently
 - this is know as “locality of reference”

Working Set

- contemporary dynamic paging algorithms are based on the idea of a “working set”
- the set of pages that a program is currently using is called the “working set”
 - if we are able to load the entire working set into memory, this process will not have any faults
 - if we are not able to load the working set, the program will thrash
- it is known that working sets tend to change fairly slowly over time, therefore we should be able to develop an algorithm that monitors the working set, and:
 - releases frames that are no longer needed (so other processes can use them)
 - acquires additional frames when the number of pages in the working set increases

Window Size

- for working set algorithms, we add a parameter w that defines a “window size”
- this window size defines the number of most recent page references that are used to define the working set, for example:
 - assume $w = 4$

t_i
↓

Request			3	2	0	5	5	3	2	1	3	1	3	2	4	5	6	5
Frame	0	-	<u>3</u>	3	3	3												
	1	-		<u>2</u>	2	2												
	2	-			<u>0</u>	0												
	3	-				<u>5</u>												
Page faults			*	*	*	*												

Working Set - 1

- in working set algorithms, new pages are added into frames that are not currently in the working set
 - is a page is replaced because it is not in the working set, it still must be written to disk
- continuing for example with $w = 4$

the window size is 4; it moves as t_i increases

the working set at this time is { 3 2 0 5 }

Request			3	2	0	5	5	3	2	1	3	1	3	2	4	5	6	5
Frame	0	-	<u>3</u>	3	3	3												
	1	-		<u>2</u>	2	2												
	2	-			<u>0</u>	0												
	3	-				<u>5</u>												
Page faults			*	*	*	*												

Working Set - 2

*now frame 0 is
available to other
processes*

the window size is 4; it
moves as t_i increases

the working set at this
time is { 2 0 5 }

t_i



Request			3	2	0	5	5	3	2	1	3	1	3	2	4	5	6	5
Frame	0	-	<u>3</u>	3	3	3												
	1	-		<u>2</u>	2	2	2											
	2	-			<u>0</u>	0	0											
	3	-				<u>5</u>	5											
Page faults			*	*	*	*												

Working Set - 3

page fault occurs when page 3 joins the working set

frame 1 leaves the working set when page 2 leaves the window

the window size is 4; it moves as t_i increases

the working set is now { 3 0 5 }

t_i

Request			3	2	0	5	5	3	2	1	3	1	3	2	4	5	6	5
Frame	0	-	<u>3</u>	3	3	3		<u>3</u>										
	1	-		<u>2</u>	2	2	2											
	2	-			<u>0</u>	0	0	0										
	3	-				<u>5</u>	5	5										
Page faults			*	*	*	*		*										

Working Set - 4

page fault occurs when page 2 joins the working set

frame 2 leaves the working set when page 0 leaves the window

the window size is 4; it moves as t_i increases

the working set at this time is { 3 2 5 }

t_i

Request			3	2	0	5	5	3	2	1	3	1	3	2	4	5	6	5
Frame	0	-	<u>3</u>	3	3	3		<u>3</u>	3									
	1	-		<u>2</u>	2	2	2		<u>2</u>									
	2	-			<u>0</u>	0	0	0										
	3	-				<u>5</u>	5	5	5									
Page faults			*	*	*	*		*	*									

Working Set - 5

- continuing the example ...

the working set at this time is { 5 3 2 1 }

t_i

Request			3	2	0	5	5	3	2	1	3	1	3	2	4	5	6	5
Frame	0	-	<u>3</u>	3	3	3		<u>3</u>	3	3								
	1	-		<u>2</u>	2	2	2		<u>2</u>	2								
	2	-			<u>0</u>	0	0	0		<u>1</u>								
	3	-				<u>5</u>	5	5	5	5								
Page faults			*	*	*	*		*	*	*								

Working Set - 6

- continuing the example ...

the working set at this time is { 3 2 1 }

t_i
↓

Request			3	2	0	5	5	3	2	1	3	1	3	2	4	5	6	5
Frame	0	-	<u>3</u>	3	3	3		<u>3</u>	3	3	3							
	1	-		<u>2</u>	2	2	2		<u>2</u>	2	2							
	2	-			<u>0</u>	0	0	0		<u>1</u>	1							
	3	-				<u>5</u>	5	5	5	5								
Page faults			*	*	*	*		*	*	*								

Working Set - 7

- continuing the example ...

the working set at this time is { 3 2 1 }

t_i
↓

Request			3	2	0	5	5	3	2	1	3	1	3	2	4	5	6	5
Frame	0	-	<u>3</u>	3	3	3		<u>3</u>	3	3	3	3						
	1	-		<u>2</u>	2	2	2		<u>2</u>	2	2	2						
	2	-			<u>0</u>	0	0	0		<u>1</u>	1	1						
	3	-				<u>5</u>	5	5	5	5								
Page faults			*	*	*	*		*	*	*								

Working Set - 8

- continuing the example ...

the working set at this time is { 3 1 }

Request			3	2	0	5	5	3	2	1	3	1	3	2	4	5	6	5
Frame	0	-	<u>3</u>	3	3	3		<u>3</u>	3	3	3	3	3					
	1	-		<u>2</u>	2	2	2		<u>2</u>	2	2	2						
	2	-			<u>0</u>	0	0	0		<u>1</u>	1	1	1					
	3	-				<u>5</u>	5	5	5	5								
Page faults			*	*	*	*		*	*	*								

t_i
↓

Working Set - 9

- continuing the example ...

the working set at this time is { 3 1 2 }

t_i



Request			3	2	0	5	5	3	2	1	3	1	3	2	4	5	6	5
Frame	0	-	<u>3</u>	3	3	3		<u>3</u>	3	3	3	3	3	3				
	1	-		<u>2</u>	2	2	2		<u>2</u>	2	2	2	2	<u>2</u>				
	2	-			<u>0</u>	0	0	0		<u>1</u>	1	1	1	1				
	3	-				<u>5</u>	5	5	5	5								
Page faults			*	*	*	*		*	*	*				*				

Working Set - 10

- continuing the example ...

the working set at this time is { 1 2 3 4 }

t_i
↓

Request			3	2	0	5	5	3	2	1	3	1	3	2	4	5	6	5
Frame	0	-	<u>3</u>	3	3	3		<u>3</u>	3	3	3	3	3	3	3			
	1	-		<u>2</u>	2	2	2		<u>2</u>	2	2	2		<u>2</u>	2			
	2	-			<u>0</u>	0	0	0		<u>1</u>	1	1	1	1	1			
	3	-				<u>5</u>	5	5	5	5					<u>4</u>			
Page faults			*	*	*	*		*	*	*				*	*			

Working Set - 11

- continuing the example ...

page 5 needs a home, but
page 1 fell out of the window,
so we can put page 5 in frame 2

the working set at this
time is { 2 3 4 5 }

t_i
↓

Request			3	2	0	5	5	3	2	1	3	1	3	2	4	5	6	5
Frame	0	-	<u>3</u>	3	3	3		<u>3</u>	3	3	3	3	<u>3</u>	3	3	3		
	1	-		<u>2</u>	2	2	2		<u>2</u>	2	2	2		<u>2</u>	2	2		
	2	-			<u>0</u>	0	0	0		<u>1</u>	1	1	1	1	1	<u>5</u>		
	3	-				<u>5</u>	5	5	5	5					<u>4</u>	4		
Page faults			*	*	*	*		*	*	*				*	*	*		

Working Set - 12

- continuing the example ...

same thing happens ...
page 3 leaves the set ...
... and page 6 joins

the working set at this
time is { 2 4 5 6 }

t_i
↓

Request			3	2	0	5	5	3	2	1	3	1	3	2	4	5	6	5
Frame	0	-	<u>3</u>	3	3	3		<u>3</u>	3	3	3	3	3	3	3	3	<u>6</u>	
	1	-		<u>2</u>	2	2	2		<u>2</u>	2	2	2		<u>2</u>	2	2	2	
	2	-			<u>0</u>	0	0	0		<u>1</u>	1	1	1	1	1	<u>5</u>	5	
	3	-				<u>5</u>	5	5	5	5					<u>4</u>	4	4	
Page faults			*	*	*	*		*	*	*			*	*	*	*		

Working Set - 13

- continuing the example ...

the working set at this time is { 4 5 6 }

t_i
↓

Request			3	2	0	5	5	3	2	1	3	1	3	2	4	5	6	5
Frame	0	-	<u>3</u>	3	3	3		<u>3</u>	3	3	3	3	3	3	3	3	<u>6</u>	6
	1	-		<u>2</u>	2	2	2		<u>2</u>	2	2	2		<u>2</u>	2	2	2	
	2	-			<u>0</u>	0	0	0		<u>1</u>	1	1	1	1	1	<u>5</u>	5	5
	3	-				<u>5</u>	5	5	5	5					<u>4</u>	4	4	4
Page faults			*	*	*	*		*	*	*			*	*	*	*		

Locality of Reference (revisited)

- If we look at the working set we see it changes slowly ...

<u>Set</u>	<u>Size</u>
{ 3 }	1
{ 3 2 }	2
{ 3 2 0 }	3
{ 3 2 0 5 }	4
{ 2 0 5 }	3
{ 3 0 5 }	3
{ 3 2 5 }	3
{ 3 2 5 1 }	4
{ 3 2 1 }	3
{ 3 2 1 }	3
{ 3 1 }	2
{ 3 2 1 }	3
{ 3 2 1 4 }	4
{ 3 2 5 4 }	4
{ 6 2 5 4 }	4
{ 6 5 4 }	3

What happens with different window sizes?

- Lets try $w = 2$
- same reference stream

t_i
↓

Request			3	2	0	5	5	3	2	1	3	1	3	2	4	5	6	5
Frame	0	-	<u>3</u>	3														
	1	-		<u>2</u>														
Page faults				*	*													

Example 2 - $\omega = 2$ (1)

t_i
↓

Request			3	2	0	5	5	3	2	1	3	1	3	2	4	5	6	5
Frame	0	-	<u>3</u>	3	<u>0</u>													
	1	-		<u>2</u>	2													
Page faults			*	*	*													

Example 2 - $\omega = 2$ (2)

t_i
↓

Request			3	2	0	5	5	3	2	1	3	1	3	2	4	5	6	5
Frame	0	-	<u>3</u>	3	<u>0</u>	0												
	1	-		<u>2</u>	2	<u>5</u>												
Page faults			*	*	*	*												

Example 2 - $\omega = 2$ (3)

- we can see that the number of page faults increases
- What is the maximum number of page frames that this process could use?

7 -> one for each frame 0-6

Request			3	2	0	5	5	3	2	1	3	1	3	2	4	5	6	5
Frame	0	-	<u>3</u>	3	<u>0</u>	0		<u>3</u>	3	<u>1</u>	1	1	1	<u>2</u>	2	<u>5</u>	5	5
	1	-		<u>2</u>	2	<u>5</u>	5	5	<u>2</u>	2	<u>3</u>	3	3	3	<u>4</u>	4	<u>6</u>	6
Page faults			*	*	*	*		*	*	*	*			*	*	*	*	

t_i
↓

Example 3 - $\omega = 7$ (1)

- OK, so lets see what happens if we set the window size to 7

t_i
↓

Request			3	2	0	5	5	3	2	1	3	1	3	2	4	5	6	5
Frame	0	-	<u>3</u>	3	3	3	3	3	3									
	1	-		<u>2</u>	2	2	2	2	2									
	2	-			<u>0</u>	0	0	0	0									
	3	-				<u>5</u>	5	5	5									
	4	-																
	5	-																
	6	-																
Page faults			*	*	*	*												

Example 3 - $\omega = 7$ (2)

- OK, so lets see what happens if we set the window size to 7

t_i
↓

Request			3	2	0	5	5	3	2	1	3	1	3	2	4	5	6	5
Frame	0	-	<u>3</u>	3	3	3	3	3	3	3								
	1	-		<u>2</u>	2	2	2	2	2	2								
	2	-			<u>0</u>	0	0	0	0	0								
	3	-				<u>5</u>	5	5	5	5								
	4	-								<u>1</u>								
	5	-																
	6	-																
Page faults			*	*	*	*				*								

Example 3 - $\omega = 7$ (3)

- OK, so lets see what happens if we set the window size to 7

t_i
↓

Request			3	2	0	5	5	3	2	1	3	1	3	2	4	5	6	5
Frame	0	-	<u>3</u>	3	3	3	3	3	3	3	3							
	1	-		<u>2</u>	2	2	2	2	2	2	2							
	2	-			<u>0</u>	0	0	0	0	0	0							
	3	-				<u>5</u>	5	5	5	5	5							
	4	-								<u>1</u>	1							
	5	-																
	6	-																
Page faults			*	*	*	*				*								

Example 3 - $\omega = 7$ (4)

- OK, so lets see what happens if we set the window size to 7

t_i
↓

Request			3	2	0	5	5	3	2	1	3	1	3	2	4	5	6	5
Frame	0	-	<u>3</u>	3	3	3	3	3	3	3	3	3						
	1	-		<u>2</u>	2	2	2	2	2	2	2	2						
	2	-			<u>0</u>	0	0	0	0	0	0							
	3	-				<u>5</u>	5	5	5	5	5	5						
	4	-								<u>1</u>	1	1						
	5	-																
	6	-																
Page faults			*	*	*	*				*								

Example 3 - $\omega = 7$ (5)

- OK, so lets see what happens if we set the window size to 7

t_i
↓

Request			3	2	0	5	5	3	2	1	3	1	3	2	4	5	6	5
Frame	0	-	<u>3</u>	3	3	3	3	3	3	3	3	3	3					
	1	-		<u>2</u>	2	2	2	2	2	2	2	2	2					
	2	-			<u>0</u>	0	0	0	0	0	0							
	3	-				<u>5</u>	5	5	5	5	5	5	5					
	4	-								<u>1</u>	1	1	1					
	5	-																
	6	-																
Page faults			*	*	*	*				*								

Example 3 - $\omega = 7$ (6)

- we see that the number of page faults has decreased
- we also observe that we don't usually require 7 frames, as the locality of reference is not that big

t_i
↓

Request			3	2	0	5	5	3	2	1	3	1	3	2	4	5	6	5
Frame	0	-	<u>3</u>	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3
	1	-		<u>2</u>	2	2	2	2	2	2	2	2	2	2	2	2	2	2
	2	-			<u>0</u>	0	0	0	0	0	0				<u>4</u>	4	4	4
	3	-				<u>5</u>	5	5	5	5	5	5	5			<u>5</u>	5	5
	4	-								<u>1</u>	1	1	1	1	1	1	1	1
	5	-															<u>6</u>	6
	6	-																
Page faults			*	*	*	*				*					*	*	*	

Clock Algorithms

- so far we have looked at working set algorithms from the perspective of a single process
- to make this type of algorithm useful, we need a way to implement it for many processes
- first some background and assumptions ...
 1. we assume that page faults are not that frequent,
 - ie, there will be many page references in many processes between page swaps
 2. we assume that whenever a page is referenced, a reference bit R is set
 3. we assume that each process keeps a current virtual time
 - this is the amount of CPU time the process has actually used since creation
 4. we assume that each page stores a “time of last reference”
 - this is the time that the page was last accessed

Simple Clock Algorithm (1)

- we define a circular linked list
- as pages are loaded, they are added to the list
- every time a page is referenced its R bit is set to 1
- when a page is loaded, its R-bit is set to 1 and all other frames R-bits are cleared (set to zero)


Frame	R-bit	Process
10	1	P3

This tells us that the page in frame 10 was accessed by Process 3

Simple Clock Algorithm (2)

- we define a circular linked list
- as pages are loaded, they are added to the list
- every time a page is referenced its R bit is set to 1
- when a page is loaded, its R-bit is set to 1 and all other frames R-bits are cleared (set to zero)


Frame	R-bit	Process
10	0	P3
4	1	P7



Simple Clock Algorithm (3)

- we define a circular linked list
- as pages are loaded, they are added to the list
- every time a page is referenced its R bit is set to 1
- when a page is loaded, its R-bit is set to 1 and all other frames R-bits are cleared (set to zero)


Frame	R-bit	Process
10	0	P3
4	0	P7
53	1	P9



Simple Clock Algorithm (4)

- we define a circular linked list
- as pages are loaded, they are added to the list
- every time a page is referenced its R bit is set to 1
- when a page is loaded, its R-bit is set to 1 and all other frames R-bits are cleared (set to zero)


Frame	R-bit	Process
10	0	P3
4	0	P7
53	0	P9
9	1	P3



Simple Clock Algorithm (5)

- we define a circular linked list
- as pages are loaded, they are added to the list
- every time a page is referenced its R bit is set to 1
- when a page is loaded, its R-bit is set to 1 and all other frames R-bits are cleared (set to zero)


Frame	R-bit	Process
10	0	P3
4	0	P7
53	0	P9
9	0	P3
34	1	P2



Simple Clock Algorithm (6)

- we define a circular linked list
- as pages are loaded, they are added to the list
- every time a page is referenced its R bit is set to 1
- when a page is loaded, its R-bit is set to 1 and all other frames R-bits are cleared (set to zero)


Frame	R-bit	Process
10	0	P3
4	0	P7
53	0	P9
9	0	P3
34	0	P2
19	1	P4



Simple Clock Algorithm (7)

- we define a circular linked list
- as pages are loaded, they are added to the list
- every time a page is referenced its R bit is set to 1
- when a page is loaded, its R-bit is set to 1 and all other frames R-bits are cleared (set to zero)

Frame	R-bit	Process
10	0	P3
4	0	P7
53	0	P9
9	0	P3
34	0	P2
19	0	P4
48	1	P4



Simple Clock Algorithm (8)

- we define a circular linked list
- as pages are loaded, they are added to the list
- every time a page is referenced its R bit is set to 1
- when a page is loaded, its R-bit is set to 1 and all other frames R-bits are cleared (set to zero)

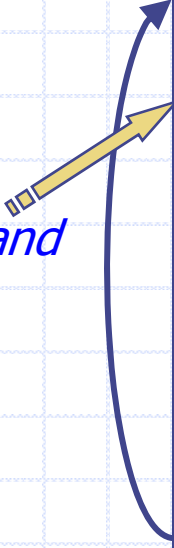
clock hand

Frame	R-bit	Process
10	0	P3
4	0	P7
53	0	P9
9	0	P3
34	0	P2
19	0	P4
48	0	P4
29	1	P3

- we also define a “clock hand” that points to pages in the circular list
- this hand will be moved by the algorithm as it looks for pages to replace

Simple Clock Algorithm (9)

- assume that no page faults happen for a while
- processes get scheduled and use pages they have loaded, setting R-bits

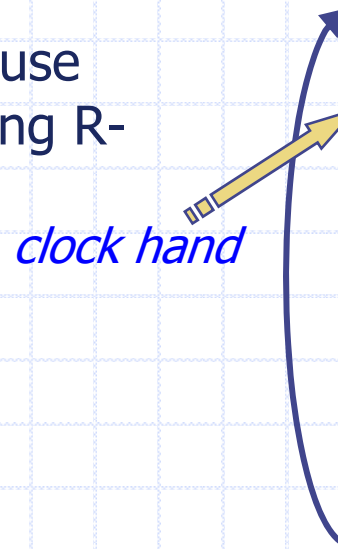


A diagram of a clock hand, labeled "clock hand" in blue italicized text, is shown pointing to the 4th row of the table (the row with Frame 53). The hand is a yellow arrow with a black outline, and a curved black arrow indicates its path from the bottom row up to the 4th row.

Frame	R-bit	Process
10	0	P3
4	0	P7
53	0	P9
9	0	P3
34	0	P2
19	0	P4
48	0	P4
29	1	P3

Simple Clock Algorithm (10)

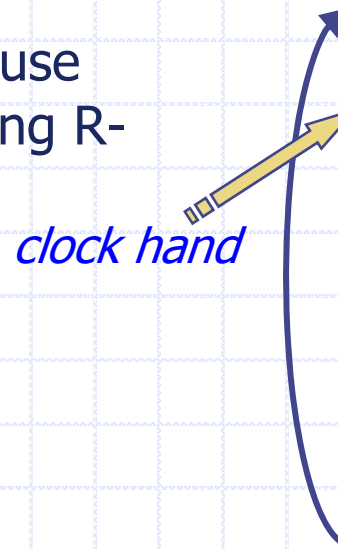
- assume that no page faults happen for a while
- processes get scheduled and use pages they have loaded, setting R-bits



Frame	R-bit	Process
10	0	P3
4	0	P7
53	1	P9
9	0	P3
34	0	P2
19	0	P4
48	0	P4
29	1	P3

Simple Clock Algorithm (11)

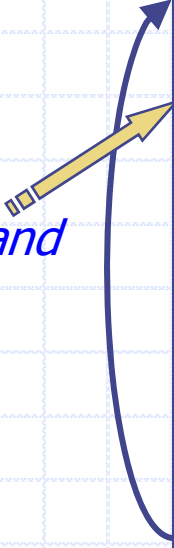
- assume that no page faults happen for a while
- processes get scheduled and use pages they have loaded, setting R-bits



Frame	R-bit	Process
10	0	P3
4	1	P7
53	1	P9
9	0	P3
34	0	P2
19	0	P4
48	0	P4
29	1	P3

Simple Clock Algorithm (12)

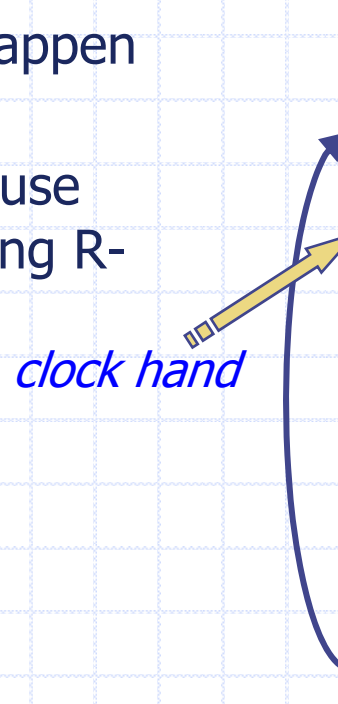
- assume that no page faults happen for a while
- processes get scheduled and use pages they have loaded, setting R-bits



Frame	R-bit	Process
10	0	P3
4	1	P7
53	1	P9
9	0	P3
34	0	P2
19	0	P4
48	1	P4
29	1	P3

Simple Clock Algorithm (13)

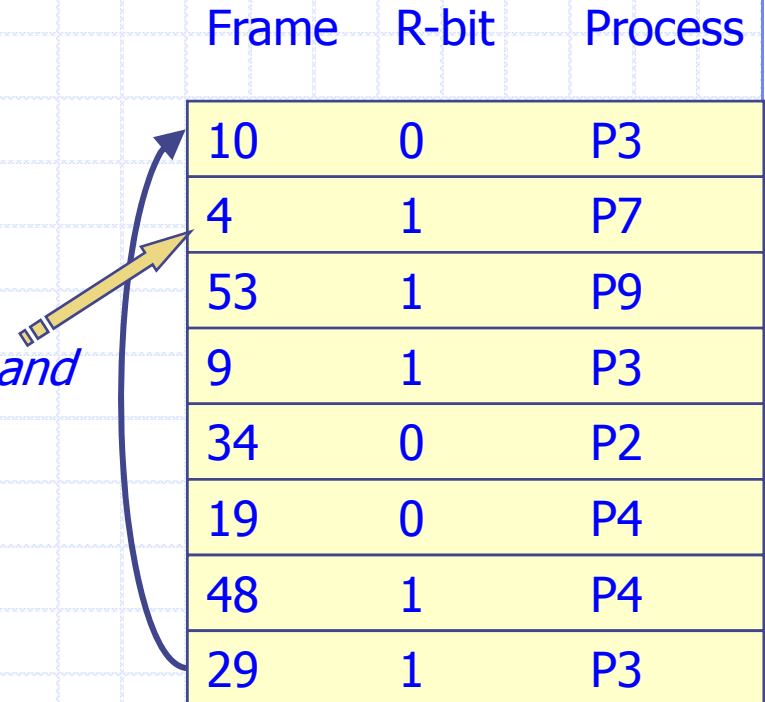
- assume that no page faults happen for a while
- processes get scheduled and use pages they have loaded, setting R-bits



Frame	R-bit	Process
10	0	P3
4	1	P7
53	1	P9
9	1	P3
34	0	P2
19	0	P4
48	1	P4
29	1	P3

Simple Clock Algorithm (14)

- at this point let us assume that process P3 causes a page fault
- the “clock hand” is pointing to frame 4
- the algorithm examines pages, starting with the frame pointed to by the clock hand
 - if the R-bit is set, the page must be in the current working set
 - the algorithm clears it and moves to the next frame
 - if the R-bit is not set, the page is not in the working set and can be replaced
 - the page is moved to disk and the frame is loaded with the requested page
 - if all R-bits are set (ie: we go all the way around without finding a page that is not in the working set)
 - we select a victim at random



Frame	R-bit	Process
10	0	P3
4	1	P7
53	1	P9
9	1	P3
34	0	P2
19	0	P4
48	1	P4
29	1	P3

Simple Clock Algorithm (15)

- Frame 4 is in the working set ...
... clear R-bit and try the next one ...

clock hand

Frame	R-bit	Process
10	0	P3
4	1	P7
53	1	P9
9	1	P3
34	0	P2
19	0	P4
48	1	P4
29	1	P3

Simple Clock Algorithm (16)

- Frame 53 is in the working set ...
... clear R-bit and try the next one ...

clock hand

Frame	R-bit	Process
10	0	P3
4	0	P7
53	1	P9
9	1	P3
34	0	P2
19	0	P4
48	1	P4
29	1	P3

Simple Clock Algorithm (17)

- Frame 9 is in the working set ...
... clear R-bit and try the next one ...


clock hand 

Frame	R-bit	Process
10	0	P3
4	0	P7
53	0	P9
9	1	P3
34	0	P2
19	0	P4
48	1	P4
29	1	P3

Simple Clock Algorithm (18)

- Frame 34 is NOT in the working set
- frame 34 is selected for replacement
- the page that is loaded in frame 34 is written to disk (if dirty bit is set)

clock hand



Frame	R-bit	Process
10	0	P3
4	0	P7
53	0	P9
9	0	P3
34	0	P2
19	0	P4
48	1	P4
29	1	P3

Simple Clock Algorithm (19)

- Frame 34 is NOT in the working set
- frame 34 is selected for replacement
- the page that is loaded in frame 34 is written to disk (if dirty bit is set)
- the new page is loaded

clock hand



Frame	R-bit	Process
10	0	P3
4	0	P7
53	0	P9
9	0	P3
34	1	P3
19	0	P4
48	1	P4
29	1	P3

Simple Clock Algorithm (20)

- Frame 34 is NOT in the working set
- frame 34 is selected for replacement
- the page that is loaded in frame 34 is written to disk (if dirty bit is set)
- the new page is loaded
- all R-bits are cleared (except for the newly loaded page), and execution continues

Frame	R-bit	Process
10	0	P3
4	0	P7
53	0	P9
9	0	P3
34	1	P3
19	0	P4
48	0	P4
29	0	P3

clock hand



WSClock Algorithm

- This is the algorithm that is used for page replacement in most dynamic paging systems.
- It is basically the same as the clock algorithm presented previously, except:
 1. uses the "time of last reference" to make sure that there is a window ω to keep pages in the working set during times of frequent paging activity
 2. to define the window ω we work with a time interval (as opposed to counting page references)
 3. when the algorithm identifies a page with R-bit == 1, it is in the working set, so we update "time of last reference", clear the R-bit, and move on
 4. when R-bit == 0, the algorithm compares last use to see if it is in the window, ie:
`if (current_virtual_time - time_of_last_reference) > ω`
... then replace the page

WSClock Example (from your textbook)

- Assume 3 processes, with current virtual times (cvt) as follows:
 - P0: cvt = 55
 - P1: cvt = 75
 - P2: cvt = 80
- Assume $\omega = 25$
- We will show each page frame entry like this:


Frame	R-bit	Process	LastRef
-------	-------	---------	---------

F: 10	R: 0	P: 3	Ref: 15
--------------	-------------	-------------	----------------

$\omega = 25$ P0: cvt = 55
P1: cvt = 75
P2: cvt = 80

WSClock Example (2)

Assume a page fault has just occurred in P0.



WSClock - check R-bit for frame 6
- page is in working set (because R=1) ...
... so set R= 0 and set ref=cvt
... advance to next frame

$\omega = 25$

P0: cvt = 55

P1: cvt = 75

P2: cvt = 80

F: 14 R: 0 P: 2 Ref: 47

F: 0 R: 0 P: 0 Ref: 15

F: 13 R: 0 P: 1 Ref: 25

F: 1 R: 1 P: 0 Ref: 51

F: 12 R: 0 P: 0 Ref: 45

F: 2 R: 0 P: 1 Ref: 69

F: 11 R: 0 P: 2 Ref: 25

F: 3 R: 1 P: 2 Ref: 65

F: 10 R: 0 P: 1 Ref: 23

F: 4 R: 1 P: 2 Ref: 80

F: 9 R: 1 P: 0 Ref: 54

F: 5 R: 0 P: 1 Ref: 15

F: 8 R: 1 P: 2 Ref: 70

F: 6 R: 1 P: 1 Ref: 65

F: 7 R: 0 P: 0 Ref: 33

Comp 4735

Page 101

WSClock - check R-bit for frame 7

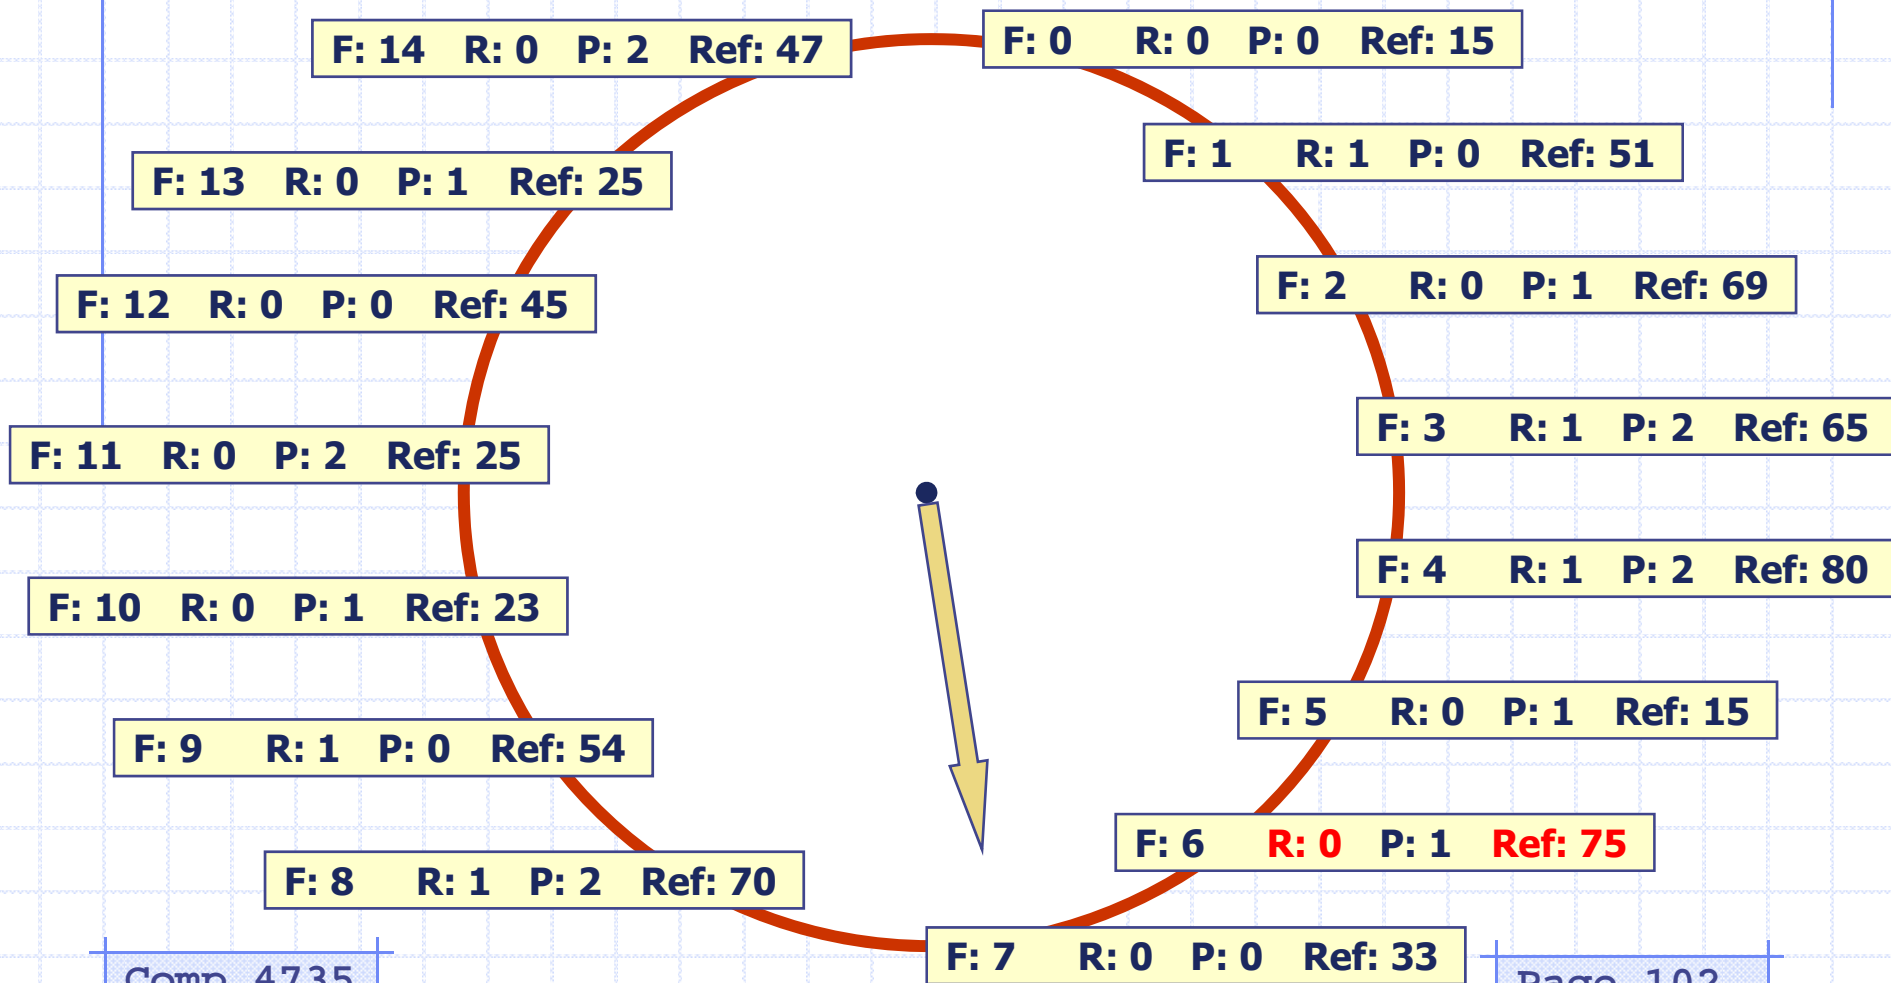
- $R = 0$, so we test $(cvt_0 - ref_0 < \omega)$
- $(55 - 33) < 25 \dots 22 < 25$
- therefore F7 is still in P0's working set
... and we advance to next frame

$\omega = 25$

P0: cvt = 55

P1: cvt = 75

P2: cvt = 80



WSClock - check R-bit for frame 8
- R-bit is set ...
... move to next frame

$\omega = 25$

P0: cvt = 55

P1: cvt = 75

P2: cvt = 80

F: 14 R: 0 P: 2 Ref: 47

F: 0 R: 0 P: 0 Ref: 15

F: 13 R: 0 P: 1 Ref: 25

F: 1 R: 1 P: 0 Ref: 51

F: 12 R: 0 P: 0 Ref: 45

F: 2 R: 0 P: 1 Ref: 69

F: 11 R: 0 P: 2 Ref: 25

F: 3 R: 1 P: 2 Ref: 65

F: 10 R: 0 P: 1 Ref: 23

F: 4 R: 1 P: 2 Ref: 80

F: 9 R: 1 P: 0 Ref: 54

F: 5 R: 0 P: 1 Ref: 15

F: 8 R: 1 P: 2 Ref: 70

F: 6 R: 0 P: 1 Ref: 75

F: 7 R: 0 P: 0 Ref: 33

Comp 4735

Page 103

WSClock - check R-bit for frame 9
- R-bit is set ...
... move to next frame

$\omega = 25$

P0: cvt = 55

P1: cvt = 75

P2: cvt = 80

F: 14 R: 0 P: 2 Ref: 47

F: 0 R: 0 P: 0 Ref: 15

F: 13 R: 0 P: 1 Ref: 25

F: 1 R: 1 P: 0 Ref: 51

F: 12 R: 0 P: 0 Ref: 45

F: 2 R: 0 P: 1 Ref: 69

F: 11 R: 0 P: 2 Ref: 25

F: 3 R: 1 P: 2 Ref: 65

F: 10 R: 0 P: 1 Ref: 23

F: 4 R: 1 P: 2 Ref: 80

F: 9 R: 1 P: 0 Ref: 54

F: 5 R: 0 P: 1 Ref: 15

F: 8 R: 0 P: 2 Ref: 80

F: 6 R: 0 P: 1 Ref: 75

F: 7 R: 0 P: 0 Ref: 33

Comp 4735

Page 104

WSClock - check R-bit for frame 10

- $R = 0$... test: $(cvt_1 - ref_1 < \omega)$...
... $(75 - 23) < 25$... $(52 < 25)$

→ therefore the page is no longer in P1 W Set

$\omega = 25$

P0: cvt = 55

P1: cvt = 75

P2: cvt = 80

F: 14 R: 0 P: 2 Ref: 47

F: 0 R: 0 P: 0 Ref: 15

F: 13 R: 0 P: 1 Ref: 25

F: 1 R: 1 P: 0 Ref: 51

F: 12 R: 0 P: 0 Ref: 45

F: 2 R: 0 P: 1 Ref: 69

F: 11 R: 0 P: 2 Ref: 25

F: 3 R: 1 P: 2 Ref: 65

F: 10 R: 0 P: 1 Ref: 23

F: 4 R: 1 P: 2 Ref: 80

F: 9 R: 0 P: 0 Ref: 55

F: 5 R: 0 P: 1 Ref: 15

F: 8 R: 0 P: 2 Ref: 80

F: 6 R: 0 P: 1 Ref: 75

Comp 4735

F: 7 R: 0 P: 0 Ref: 33

Page 105

WSClock *the new page is loaded in frame 10*

- this frame is now allocated to Process 0

$\omega = 25$

P0: cvt = 55

P1: cvt = 75

P2: cvt = 80

F: 14 R: 0 P: 2 Ref: 47

F: 0 R: 0 P: 0 Ref: 15

F: 13 R: 0 P: 1 Ref: 25

F: 1 R: 1 P: 0 Ref: 51

F: 12 R: 0 P: 0 Ref: 45

F: 2 R: 0 P: 1 Ref: 69

F: 11 R: 0 P: 2 Ref: 25

F: 3 R: 1 P: 2 Ref: 65

F: 10 R: 1 P: 0 Ref: 55

F: 4 R: 1 P: 2 Ref: 80

F: 9 R: 0 P: 0 Ref: 55

F: 5 R: 0 P: 1 Ref: 15

F: 8 R: 0 P: 2 Ref: 80

F: 6 R: 0 P: 1 Ref: 75

F: 7 R: 0 P: 0 Ref: 33

Comp 4735

Page 106

WSClock

- all R-bit are reset (except the current page), and

... execution of P0 continues

$\omega = 25$

P0: cvt = 55

P1: cvt = 75

P2: cvt = 80

F: 14 R: 0 P: 2 Ref: 47

F: 0 R: 0 P: 0 Ref: 15

F: 13 R: 0 P: 1 Ref: 25

F: 1 R: 0 P: 0 Ref: 51

F: 12 R: 0 P: 0 Ref: 45

F: 2 R: 0 P: 1 Ref: 69

F: 11 R: 0 P: 2 Ref: 25

F: 3 R: 0 P: 2 Ref: 65

F: 10 R: 1 P: 0 Ref: 55

F: 4 R: 0 P: 2 Ref: 80

F: 9 R: 0 P: 0 Ref: 55

F: 5 R: 0 P: 1 Ref: 15

F: 8 R: 0 P: 2 Ref: 80

F: 6 R: 0 P: 1 Ref: 75

F: 7 R: 0 P: 0 Ref: 33

Comp 4735

Page 107

WSClock Special Cases

- what happens if the clock hand goes all the way around?
 - this means that all pages are in the working set
 - randomly select a victim and write to disk / claim it
- in practice, after a page is selected for replacement, we check if it is dirty
 - if it is dirty we schedule it to be written to disk, but don't claim it
 - we move on to the next frame, and keep scheduling dirty pages for write to disk
 - if we find a page that is not in the working set, and is clean, we claim it
 - if we don't find a clean page the hand just keeps moving; it will find one as soon as the write to disk is finished

The End