

POSIX Threads (pthreads)

- **Pthreads** is a new standard **POSIX API** that provides library routines for multithreading on **UNIX**.
- The **POSIX API** and the **Solaris API** share identical library calls, the only difference is the name of the function call.
- The **POSIX API** uses the prefix **pthread_** before each call, **Solaris API** uses the prefix **thr_** or **lwp_** before each call.
- The functionality of the two APIs however is different in several ways. The table below illustrates the similarities and differences between the two APIs.

<i>POSIX API</i>	<i>Similar Functionality</i>	<i>Solaris API</i>
Thread cancellation Scheduling policies Synchronization attributes Thread attributes	Join Key creation Create, Exit, Kill Priorities, sigmask, Mutex variables Condition variables Thread specific data	Continue Suspend Semaphore variables Concurrency setting Reader/Writer variables Daemon threads

- **Solaris** threads provide the following features not found in **pthreads**:
 - Reader/Writer locks (classic one server many client problem).
 - Ability to suspend and continue a given thread.
 - Ability to create daemon threads (**THR_DAEMON**).
 - Ability to set and get levels of concurrency.
- **pthreads** provide the following features not found in **Solaris**:
 - Ability to cancel a given thread.
 - Ability to set thread and synchronization attributes.
 - Scheduling policies.

Creating a Thread

- A new thread is created using the **pthread_create()** function. The syntax is as follows:

```
#include <pthread.h>
```

```
int pthread_create (pthread_t *thread, const pthread_attr_t *attr, void *  
(*start_routine) (void *), void *arg);
```

- The function creates a new thread with attributes specified by **attr**, within a process. If **attr** is **NULL**, default thread attributes are used.
- If the function is successful, the new **thread ID** is stored in the location referenced by **thread**.
- The third argument, **start_routine**, specifies the function which the thread is to execute. This is always a function that takes **void*** as its argument and returns **void***.
- The last parameter specifies the argument to the executing function which is **void**.
- If the function is successful, the function returns zero. Otherwise an error number is returned to identify the reason for not creating a new thread:

EAGAIN: The system lacked the resources to create another thread, or creation of a new thread will result in it exceeding the thread limit set by **PTHREAD_THREADS_MAX**.

EINVAL: The value specified by **attr** is invalid.

- Another function that is often used as part of thread creation is **pthread_join()**.

```
#include <pthread.h>
```

```
int pthread_join (pthread_t thread, void **value_ptr);
```

- This function suspends execution of the calling thread until the target function specified by **thread** terminates.
- The argument ****value_ptr** will contain the value passed to **pthread_exit()** by the terminating thread. If **pthread_exit()** is not being used then this argument is set to **NULL**.

- If the function is successful, the function returns zero. Otherwise an error number is returned to identify the reason for the failure:

ESRCH: No thread could be found corresponding to that specified by the given **thread ID**.

EDEADLK: A deadlock was detected or the value of thread specifies the calling thread.

Attribute Objects

- The pthreads API allows a programmer to create a thread in one of many **states**.
- Examples of this would be a thread that is created bound or unbound, a mutex variable that can be interprocess or intraprocess.
- Attribute objects are created and then used as arguments to creation or initialization functions to define the thread state or synchronization variables.
- The attribute objects contain all the **state information** we require to define the state of a new thread.
- The attribute object is specified when the thread is created. This is done by setting up a threads attributes structure of type **pthread_attr_t**.
- A threads attributes structure is allocated by the program and then initialized by calling:

```
int pthread_attr_init (pthread_attr_t *attr);
```

- The above call initializes the attributes structure's thread attributes to the systems default values.
- Individual attributes can then be set up by calling functions that get and set up specific attributes.
- Once an object has been set up, it can be used by any **pthread_create()** call to define the state of a new thread.
- At the end of the program the attributes structure is deinitialized by calling:

```
int pthread_attr_destroy (pthread_attr_t *attr);
```

- The state information is as follows:

Detached State:

- This state determines whether the thread will be joinable from another thread.
- When an application does not require the thread's completion status and doesn't need to know when the thread has exited, the application creates a **detached** thread.
- Detached threads cannot be the target of **pthread_join()**.
- The **detachstate** attribute specifies whether the thread is to be detached and it has one of two values:

PTHREAD_CREATE_DETACHED
PTHREAD_CREATE_JOINABLE (default value)

- This attribute is manipulated by the following two function calls:

```
// set the value of detachstate as specified in attr
int pthread_attr_setdetachstate (pthread_attr_t *attr, int detachstate);

// get the value of detachstate
int pthread_attr_getdetachstate (pthread_attr_t *attr, int detachstatep);
```

- The following code fragment illustrates how to create a detached thread:

```
int main (void)
{
    pthread_attr_t detached_attr;
    pthread_t t;
    void *detached_func (void *);

    pthread_attr_init (&detached_attr);
    pthread_attr_setdetachstate (&detached_attr,
                                PTHREAD_CREATE_DETACHED);
    pthread_create (&t, &detached_attr, detached_func, NULL);
    .....
    .....
    return(0);
}
```

Scope

- This attribute determines whether the thread has a **process-wide** or **system-wide** scope.
- The attribute can take one of the two values:

PTHREAD_SCOPE_SYSTEM **//system-wide, kernel sees the thread**

PTHREAD_SCOPE_PROCESS **//process-wide only, kernel does not see the**
 // thread - default value.

- This attribute is manipulated by the following two function calls:

```
// set the value of scopestate as specified in attr
int pthread_attr_setscope(pthread_attr_t *attr, int scopestate);
```

```
// get the value of scopestate
int pthread_attr_getscope(pthread_attr_t *attr, int scopestatep);
```

Stack Address

- Specifies the base (starting) address of the stack for the thread.
- Some care must be used when using this attribute to set specific addresses simply because this attribute cannot be used later to create new threads. It requires a change of the stack base.
- The base stack can be set and retrieved using the following functions:

pthread_attr_setstackaddr()

pthread_attr_getstackaddr()

- The default value for a stack address is **NULL**, which means that the system will assign the stack base address.

Stack Size

- This attribute specifies the size of the stack in bytes, for a thread.
- Usually a NULL size value will allow the system to set the default stack size.
- The stack size can be set and retrieved using the following functions:

pthread_attr_setstacksize()

pthread_attr_getstacksize()

Scheduling Policy

- The scheduling parameters in each thread **attribute object** define how the thread is **scheduled** to run, and set the **priority** for the thread.
- The parameters used by the scheduling policy are set and retrieved via the **schedparam** attribute in the threads attribute structure.
- This attributes object consists of a **sched_param** structure with at least one defined member:

```
struct sched_param
{
    .....
    int sched_priority;
    .....
}
```

- The priority of the thread can be set and retrieved from the **schedparam** attribute by calling:

```
// set the thread priority value to sched_priority as specified in sched_param
int pthread_attr_setschedparam (pthread_attr_t *attr, const struct
sched_param *paramp);
```

```
// get the thread priority value from sched_param
int pthread_attr_getschedparam (pthread_attr_t *attr, const struct
sched_param *paramp);
```

- In the implementation model, each thread has an integer priority (**sched_priority** in the **sched_param**) that ranges between a defined **minimum** and a defined **maximum**.
- The minimum and maximum values are defined by the kernel implementation of threads.

- The implementation's minimum and maximum priority for each **scheduling policy** can be retrieved by calling:

```
#include <sched.h>
```

```
// get minimum priority for policy
int sched_get_priority_min (int policy);
```

```
// get maximum priority for policy
int sched_get_priority_max (int policy);
```

- Where **policy** is one of the three supported policies (**SCHED_FIFO**, **SCHED_RR**, or **SCHED_OTHER**), to be discussed next.
- The threads scheduling policy is set when it is created via the **schedpolicy** attribute in the thread creation attributes (as shown earlier).
- The **schedpolicy** is set and retrieved by calling:

```
// set the thread scheduling policy to the defined integer policy
int pthread_attr_setschedpolicy (pthread_attr_t *attr, int policy);
```

```
// get the thread scheduling policy
int pthread_attr_getschedpolicy (pthread_attr_t *attr, int policy);
```

- As mentioned earlier, **POSIX** defines three thread scheduling policies:
- **SCHED_FIFO**
 - Defines a simple **first in, first out** scheduler. Threads are placed on the run queue as described below.
 - When an **active** thread is **preempted**, it is placed at the head of the queue associated with its priority.
 - When a **blocked** thread becomes **runnable**, it is placed at the tail of the queue associated with its priority.
 - When an **active** thread is the target of **pthread_setschedparam()** (set parameters after the thread is running), it is placed at the tail of the queue associated with its **new priority**.
 - When an **active** thread is the target of **sched_yield()** (force running process to relinquish CPU until process gets to the head of its run queue), it is placed at the tail of the queue associated with its priority.

- **SCHED_RR**
 - Defines a **round-robin** scheduling algorithm.
 - It is similar to the **SCHED_FIFO** except that an active **SCHED_RR** thread will be automatically preempted after it has been running for a **time quantum** associated with the entire process.
 - Following the expiration of the time quantum value, the thread is put back to the tail of its run queue.
 - The value of the time quantum can be retrieved by calling:


```
#include <sched.h>

int sched_rr_get_interval (pid_t pid, struct timespec *quantump);
```
- **SCHED_OTHER**
 - This is the same **default** policy that is also supported by Solaris. The kernel simply assigns CPU time slices based on the thread's priority.
 - In general, **SCHED_OTHER** should be used unless there are compelling reasons for strict priority scheduling, such as in real-time applications.

Synchronization Variables

- Both POSIX and Solaris threads provide several thread **synchronizing** functions, including **Semaphores** and **Mutexes**. We will discuss the simplest of these, **Mutexes**.
- **Mutexes** provide a "**locking**" mechanism that will allow only one thread at a time to execute a section of code.
- Mutexes are allocated and defined as a static or automatic data structure of type **pthread_mutex_t** as in:

```
pthread_mutex_t lock;
```

- They can also be dynamically allocated as in:

```
pthread_mutex_t *lock;
```

```
lock = (pthread_mutex_t *) malloc(sizeof (pthread_mutex_t));
```

- Once created, a mutex must be initialized as follows:

```
int pthread_mutex_init (pthread_mutex_t *lock,
    const pthread_mutexattr_t *attr);
```


- The **attr** argument points to a data structure containing attributes to be used for the mutex.
- If **attr** is **NULL**, default attributes are used by the kernel.
- Mutexes can also be statically initialized by setting them to a special value, as follows:

```
pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;
```

- This allows the mutex to be initialized without additional programming overhead.
- Mutexes are destroyed by calling:

```
pthread_mutex_destroy (pthread_mutex_t *lock);
```

- Locking and unlocking are the primary mutex operations. A mutex is locked using:

```
int pthread_mutex_lock (pthread_mutex_t *lock);
```

- The thread that **locks** the **mutex** is considered its **owner**. Only the mutex owner can unlock the mutex.

- A mutex is unlocked using:

```
int pthread_mutex_unlock (pthread_mutex_t *lock);
```

- The example program given is a simple but complete example of a multi-threaded program.
- You will need to compile it with the thread library as follows:

```
gcc multi.c -o multi -lpthread
```