

Chapter 2: Application Layer

- ❑ Principles of network applications
- ❑ Unix processes and IPC
- ❑ Web and HTTP
- ❑ Electronic Mail - SMTP, POP3, IMAP
- ❑ DNS
- ❑ Socket programming

Application Architectures

- ❑ Client-server
- ❑ Peer-to-peer (P2P)
- ❑ Hybrid of client-server and P2P

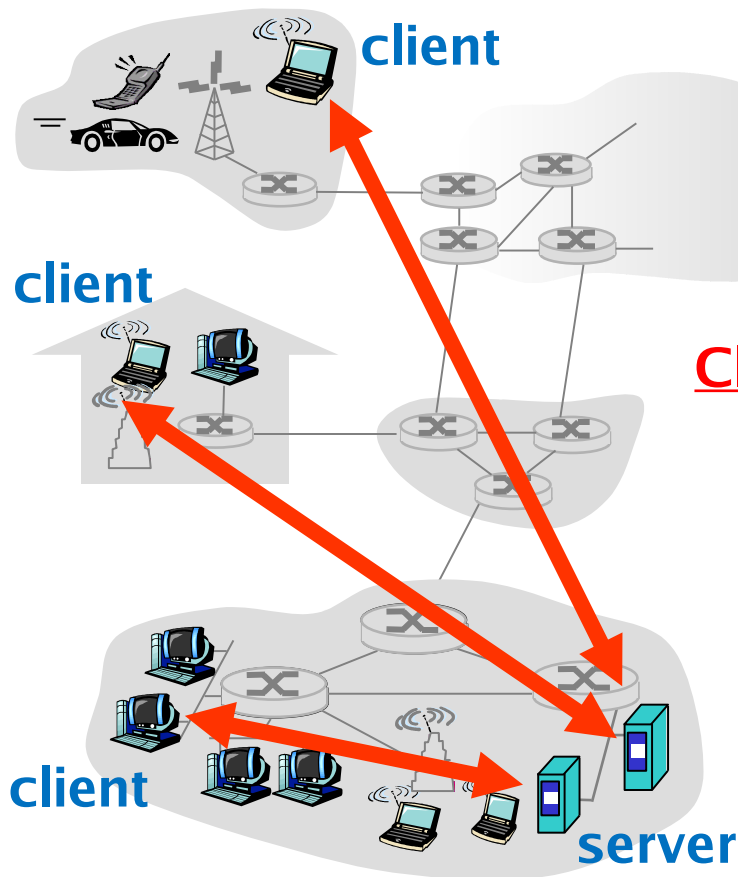
Client-server Architecture

Servers:

- ❖ Must always be online
- ❖ Provide network services to clients
- ❖ Assigned a **permanent IP address**
- ❖ **Server farms** used for scalability

Clients:


- ❖ Communicate with servers
- ❖ May be intermittently connected
- ❖ May have **dynamic IP addresses**
- ❖ Do not communicate directly with each other



Pure P2P Architecture

- ❑ No “always-on server”
- ❑ Arbitrary end systems communicate directly
- ❑ Peers are intermittently connected and change IP addresses
- ❑ example: Gnutella

Highly scalable but difficult to manage



Hybrid of Client-Server and P2P

Skype

- ❖ Voice-over-IP P2P application
- ❖ Centralized server: finding address of remote party:
- ❖ Client-client connection: direct

Instant messaging

- ❖ Chatting between two users is P2P
- ❖ Centralized service: client presence detection/location
 - User registers its IP address with central server when it comes online
 - User contacts central server to find IP addresses of buddies

Processes & Network Communications

Process: instance of an executing program.

- ❑ Processes within same host communicate using **Inter-Process Communication (IPC)**
- ❑ Processes in different hosts communicate using **sockets** (covered in TCP/IP programming)

Client process initiates communication
with a remote server

Server process accepts and processes
client requests

The UNIX Processes

- **Program** : a collection of instructions
- **Process** (= task): an instance of a program that consists 3 segments:
 - ❖ Instruction segment
 - ❖ User data segment
 - ❖ System data segment

Example

cat file1 file2

ls | wc -l

The UNIX Processes (cont.)

- ❑ **init**: a single controlling process at the top of the process tree
- ❑ System calls for process creation and manipulation:
 - ❖ **fork()**
 - ❖ **exec()**
 - ❖ **exit()**

Creating Processes

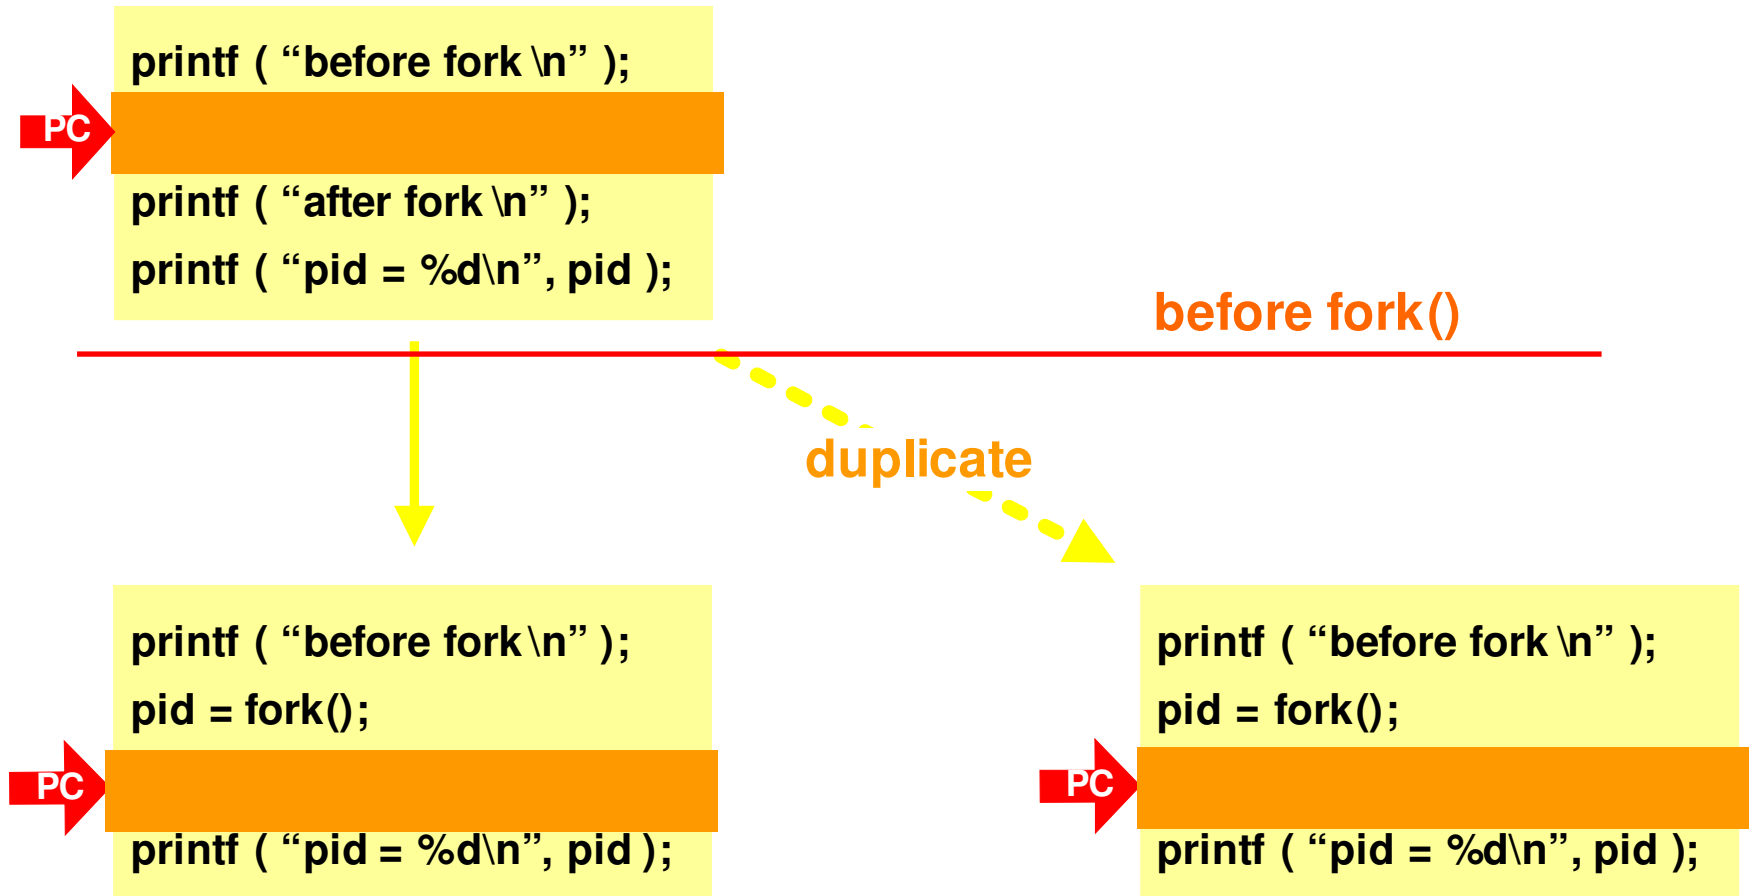
- ❑ **fork()** system call creates a new process (child) duplicating the calling process (parent)
- ❑ Parent and children run **concurrently** without synchronization

Usage

```
int pid;  
pid = fork();
```

- pid in parent = child's pid
- pid in child = 0
- -1 on failure

fork() System Call



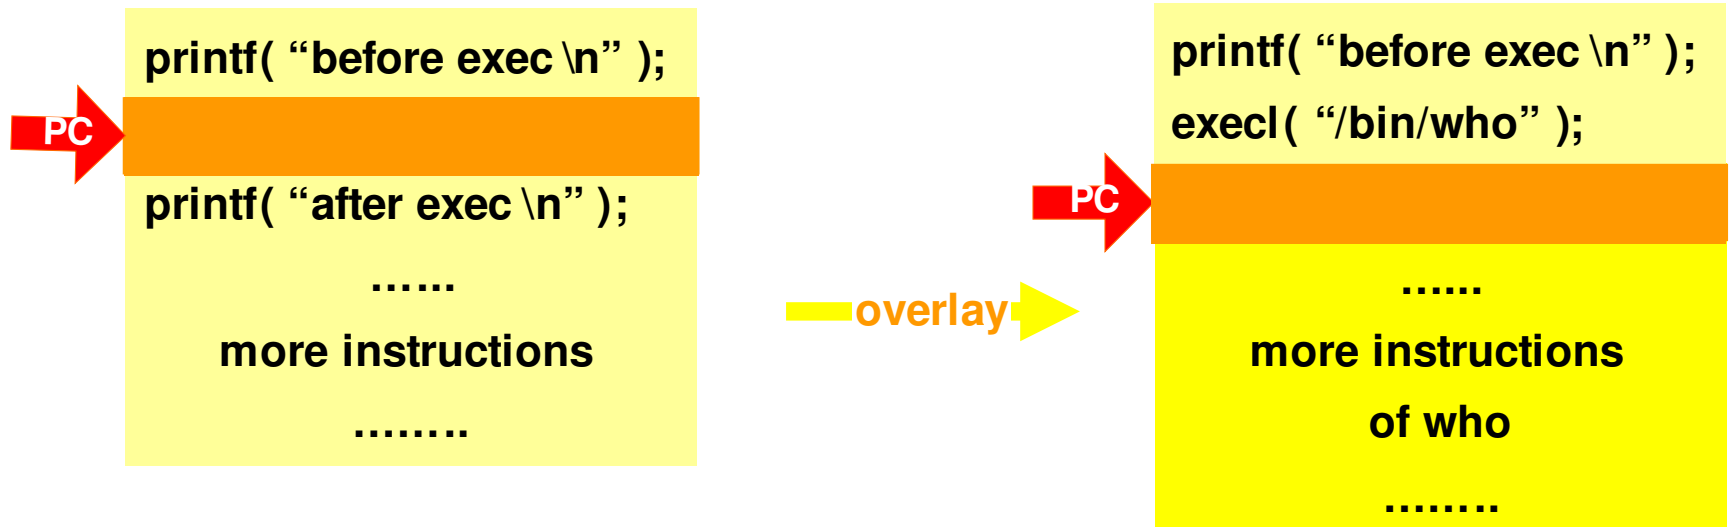
exec() System Call

- ❑ Calling process is overlaid by new code and execution begins from the first line
- ❑ No return from a successful call to exec()

```
char *path, *arg0, *arg1, ... *argn;  
execl( path, arg0, arg1, .. argn,  
      (char *)0 );
```

- ❑ **path**: path to the program executable
- ❑ **arg0**: program name
- ❑ **arg1 – n**: program parameters
- ❑ **(char *)0**: null argument list terminator
- ❑ Returns -1 for error

exec() System Call (cont.)



exit() System Call

- ❑ Terminates a process
 - ❖ Closes all open file descriptors
 - ❖ Blocked parent process is started
 - ❖ 'Exit status' indicates the success or failure
 - 0 = normal termination, non-zero = error

Usage

```
int status;  
void exit( status );
```

IPC using Pipes

- ❑ **One-way (half-duplex)** communications channel
- ❑ **Couples** one process to another
connecting the standard output of one process to the standard input of another (parent and children)
- ❑ Generalization of the UNIX **file concept**

Command example

who | wc -l

pipe() System Call

- ❑ Creates 2 file descriptors on success: **write** & **read**
- ❑ Manipulated with **read()** and **write()**
- ❑ **FIFO** data access
- ❑ Used in conjunction with **'fork'**
 - ❖ a child process will inherit any open file descriptors from the parent

pipe() System Call (cont.)

```
int filedes[2];
```

```
int retval;
```

```
retval = pipe( filedes );
```

❑ **filedes:** 2-integer array to hold the file descriptors created

- ❖ array[0] = read

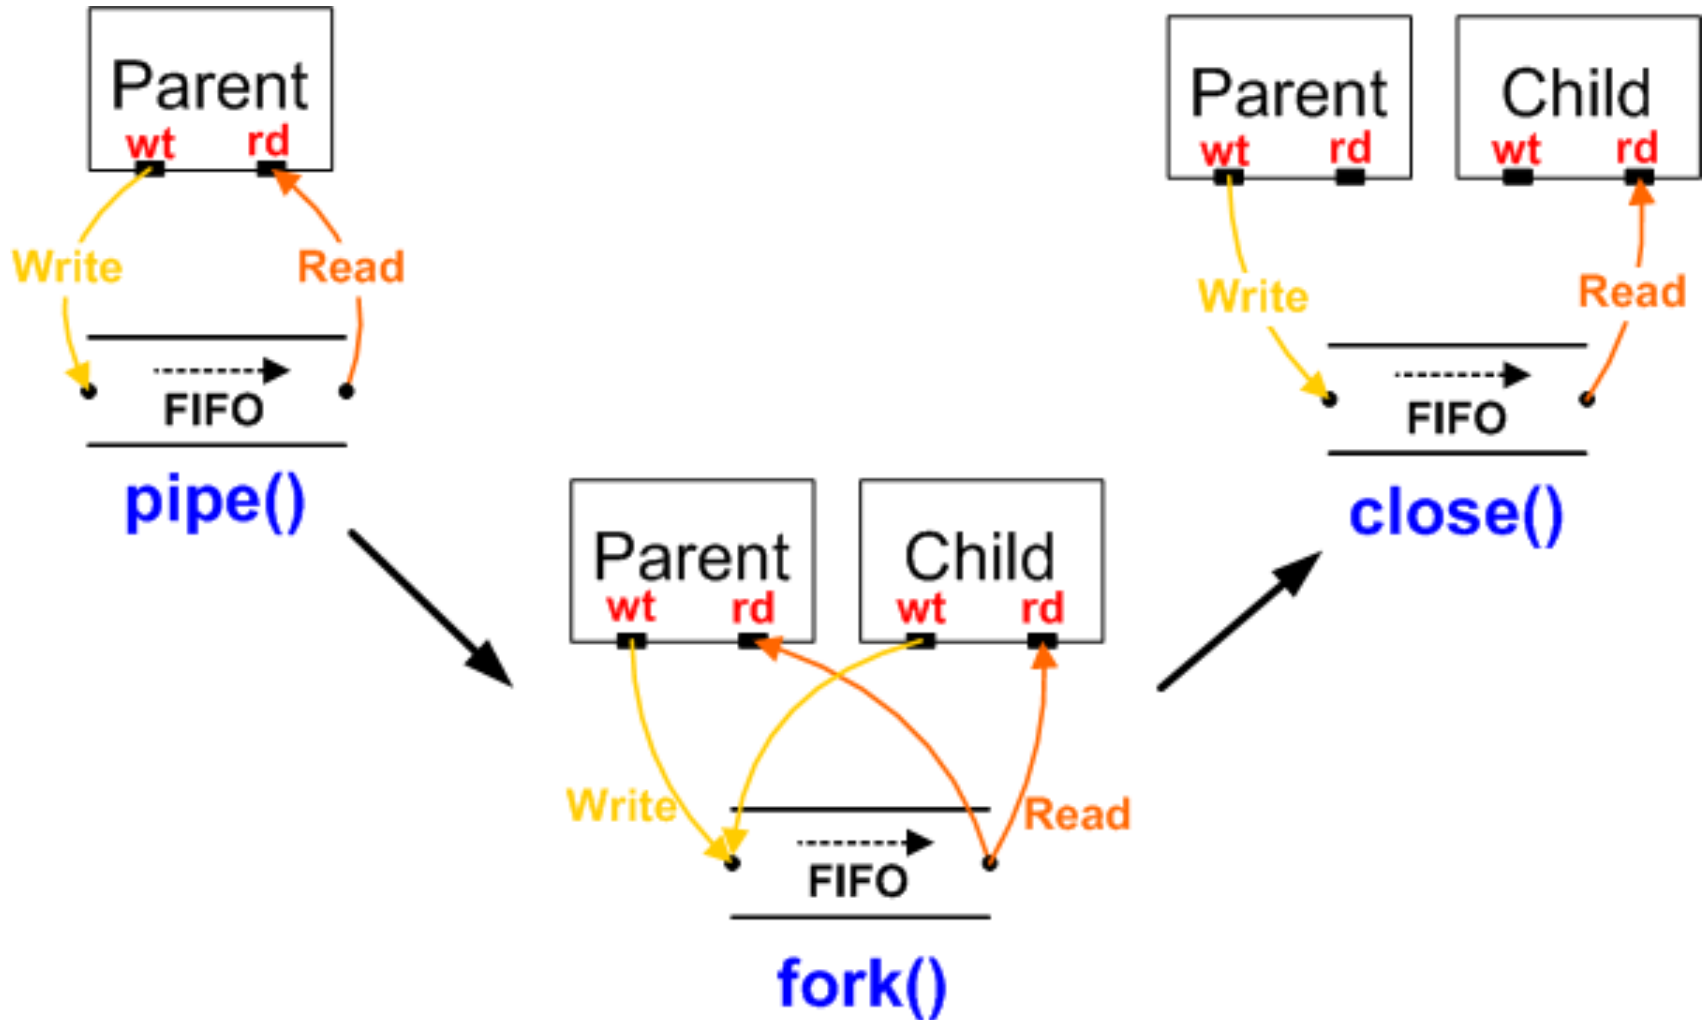
- ❖ array[1] = write

❑ **Return:**

- ❖ On success = 0.

- ❖ On error = -1, and errno is set

Pipe with Two processes



Programming with Pipes

Size : finite (5120 bytes)

- ❖ write on a pipe without enough space
 - ⇒ (default) Suspended until room is made
- ❖ read from an empty pipe
 - ⇒ (default) Blocked until data is written

Programming with Pipes (cont.)

❑ Closing write descriptor

- ❖ Other processes still have the pipe open for writing

- ⇒ No affect

- ❖ No other processes have the pipe open for writing

- ⇒ Write descriptor will be closed

- ⇒ Releases blocked reading processes if any

Programming with Pipes (cont.)

❑ Closing read descriptor

- ❖ Other processes still have the pipe open for reading

- ⇒ No affect

- ❖ No other processes have the pipe open for writing

- ⇒ Reads descriptor will be closed

- ⇒ Sends **SIGPIPE** to all writing & waiting to write processes

fork() & pipe() Example - pipe.c

- ❑ `gcc -o pipe pipe.c`
- ❑ `./pipe 3`

Application Service Requirements

Data loss

- ❑ some apps (e.g., audio) can tolerate some loss
- ❑ other apps (e.g., file transfer, ssh) require 100% reliable data transfer

Timing

- ❑ some apps (e.g., Internet telephony, interactive games) require low delay to be “effective”

Bandwidth

- ❑ some apps (e.g., multimedia) require minimum amount of bandwidth to be “effective”
- ❑ other apps (“elastic apps”) make use of whatever bandwidth they get

Transport Service Requirements For Common Applications

Application	Data loss	Bandwidth	Time Sensitive
file transfer	no loss	elastic	no
e-mail	no loss	elastic	no
Web documents	no loss	elastic	no
real-time audio/video	loss-tolerant	audio: 5kbps-1Mbps video:10kbps-5Mbps	yes, 100's msec
stored audio/video	loss-tolerant	same as above	yes, few secs
interactive games	loss-tolerant	few kbps up	yes, 100's msec
instant messaging	no loss	elastic	yes and no

Internet Transport Protocols and Services

TCP service:

- ❑ Connection-oriented: connection establishment required
- ❑ Reliable transport
- ❑ Flow control: sender prevented from overwhelming receiver
- ❑ Congestion control: throttle sender traffic when network is overloaded
- ❑ Does not provide: timing, minimum bandwidth guarantees

UDP service:

- ❑ Connectionless
- ❑ Unreliable
- ❑ Does not provide: connection setup, reliability, flow control, congestion control, timing, or bandwidth guarantee

Internet Applications: Application, Transport Protocols

<u>Application</u>	<u>Application layer protocol</u>	<u>Underlying transport protocol</u>
e-mail	SMTP [RFC 2821]	TCP
remote terminal access	Telnet [RFC 854]	TCP
Web	HTTP [RFC 2616]	TCP
file transfer	FTP [RFC 959]	TCP
streaming multimedia	proprietary (e.g. RealNetworks)	TCP or UDP
Internet telephony	proprietary (e.g., Vonage, Dialpad)	typically UDP


HTTP overview

HTTP: hypertext transfer protocol

- Web's application layer protocol
- Uses **TCP**:
- **Client** initiates TCP connection to server **port 80**
- **Server** accepts TCP connection from client

HTTP is “stateless”

- Server maintains no information about past client requests



HTTP Connections

Nonpersistent HTTP

- ❑ At most one object is sent over a TCP connection.
- ❑ HTTP/1.0 uses nonpersistent HTTP

Persistent HTTP

- ❑ Multiple objects can be sent over single TCP connection between client and server.
- ❑ HTTP/1.1 uses persistent connections in default mode

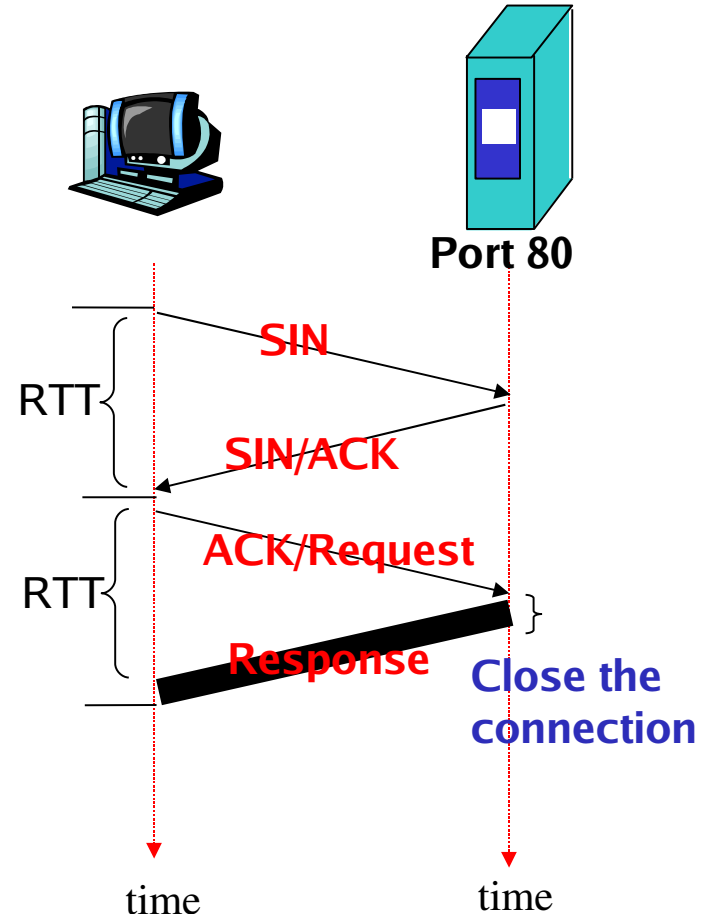
Nonpersistent HTTP

Response time:

- ❑ One RTT to initiate TCP connection
 - ❑ One RTT for HTTP request and first few bytes of HTTP response to return
 - ❑ file transmission time
- total = 2RTT + transmit time**

Nonpersistent HTTP issues:

- ❑ Requires 2 RTTs per object
- ❑ OS overhead for *each* TCP connection



Persistent HTTP

- ❑ Server leaves connection open after sending response
- ❑ Subsequent HTTP messages between same client/server sent over open connection

Persistent *without* pipelining:

- ❑ Client issues new request only when previous response has been received
- ❑ One RTT for each referenced object

Persistent *with* pipelining:

- ❑ Default in HTTP/1.1
- ❑ Client sends requests as soon as it encounters a referenced object
- ❑ As little as one RTT for all the referenced objects

User-server state: cookies

Many major Web sites use cookies

Four components:

- 1) Cookie header line of **HTTP *response* message**
- 2) Cookie header line in **HTTP *request* message**
- 3) **Cookie file** kept on user's host, managed by user's browser
- 4) **Back-end database** at Web site

Cookies can provide:

- ☐ Authorization
- ☐ Shopping carts
- ☐ Recommendations
- ☐ User session state (Web e-mail)

Cookies and privacy:

- ☐ cookies permit sites to learn a lot about you
- ☐ you may supply name and e-mail to sites

Cookies (cont.)

client

server



http request

http response
Set-cookie: 1678

Amazon server
creates ID
1678 for user


cookie file
amazon 1678

http request
cookie: 1678

http response msg

cookie-
specific
action

Access
1678



backend
database

Access
1678

one week later



http request
cookie: 1678

http response msg


cookie-
specific
action

cookie file
amazon 1678

Web caches (proxy server)

Goal: satisfy client request without involving origin server.
=> Reduce response time and traffic.

- ❑ User sets browser:
Web accesses via cache
- ❑ Browser sends all HTTP requests to cache (proxy server)
 - ❖ Object in cache: cache returns object (Hit!)
 - ❖ Else cache requests object from origin server, then returns object to client



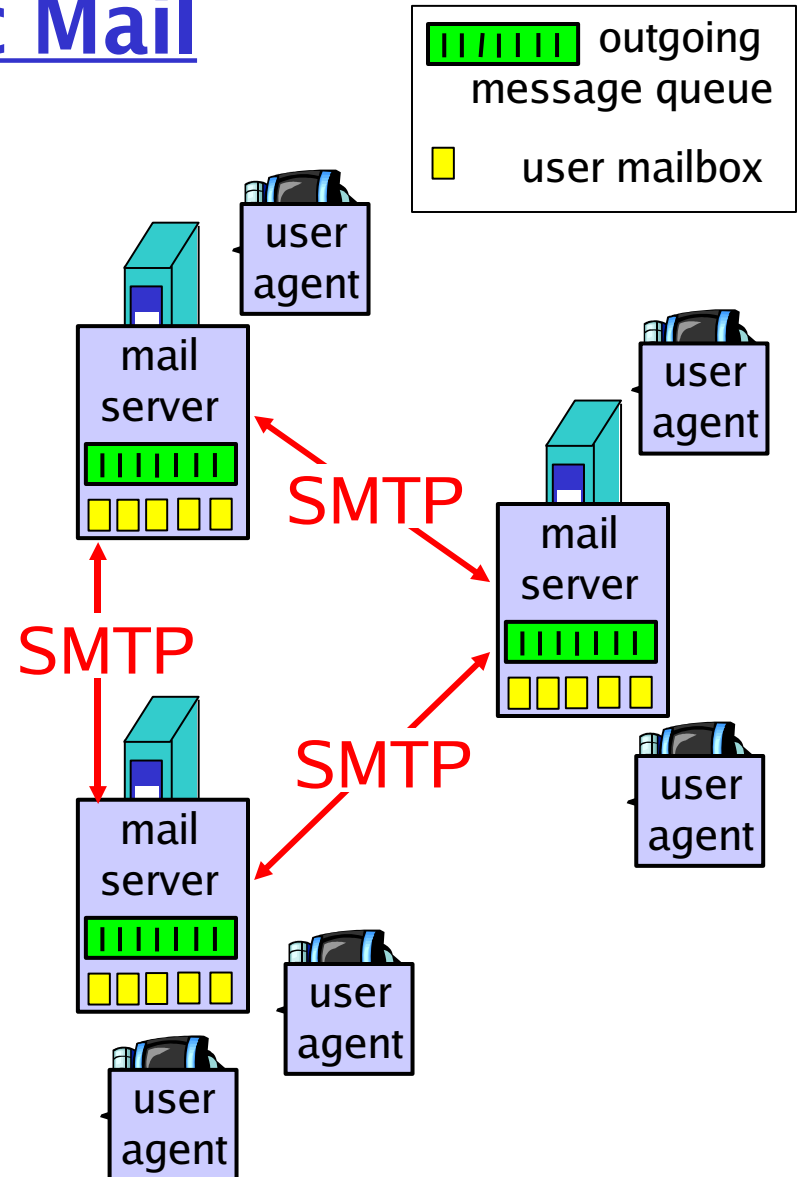
Electronic Mail

Three major components:

- ❑ User agents
- ❑ Mail servers
- ❑ Simple mail transfer protocol: **SMTP**

User Agent


- ❑ a.k.a. “email clients”
- ❑ Composing, editing, reading mail messages
- ❑ e.g., Eudora, Outlook, elm, Mozilla Thunderbird



Electronic Mail (cont.)

Mail Servers

- ❑ **Mailbox** contains incoming messages for user
- ❑ **Message queue** of outgoing (to be sent) mail messages
- ❑ **SMTP protocol** between mail servers to send email messages
 - ❖ client: sending mail server
 - ❖ “server”: receiving mail server

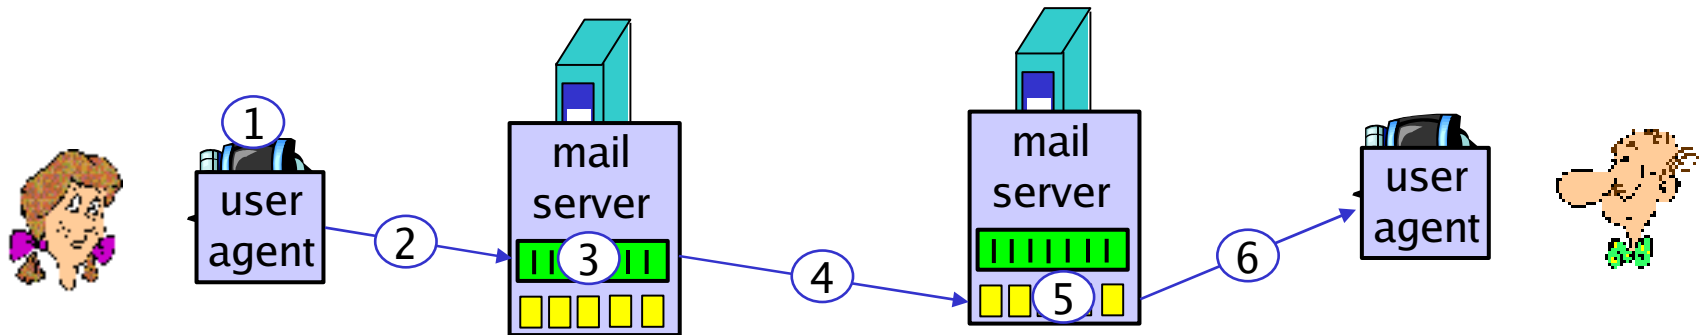


Electronic Mail: SMTP

- ❑ Uses **TCP** to reliably transfer email message from client to server, **port 25**
- ❑ Direct transfer: sending server to receiving server
- ❑ 3 phases of transfer
 - ❖ Handshaking (greeting)
 - ❖ Transfer of messages
 - ❖ Closure
- ❑ Command/response interaction
 - ❖ **Commands**: ASCII text
 - ❖ **Response**: status code and phrase
- ❑ Messages must be in 7-bit ASCII

Scenario: Alice Sends Message to Bob

- 1) Alice uses UA to compose message to bob@some school.edu
- 2) Alice's UA sends message to her mail server; message placed in message queue
- 3) Client side of SMTP opens TCP connection with Bob's mail server
- 4) SMTP client sends Alice's message over the TCP connection
- 5) Bob's mail server places the message in Bob's mailbox
- 6) Bob invokes his user agent to read message



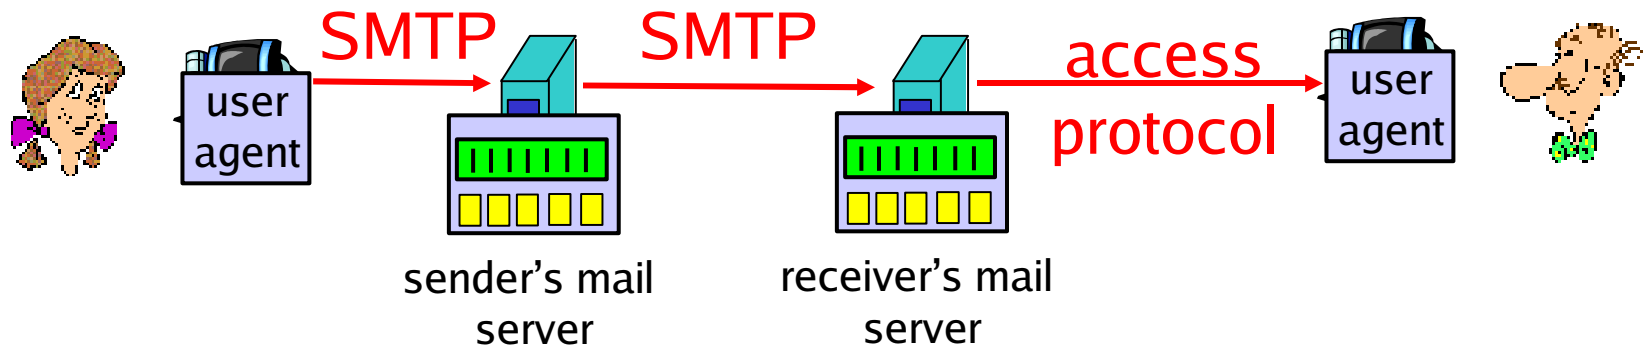
Sample SMTP Interaction

S: 220 hamburger.edu
C: HELO crepes.fr
S: 250 Hello crepes.fr, pleased to meet you
C: MAIL FROM: <alice@crepes.fr>
S: 250 alice@crepes.fr... Sender ok
C: RCPT TO: <bob@hamburger.edu>
S: 250 bob@hamburger.edu ... Recipient ok
C: DATA
S: 354 Enter mail, end with "." on a line by itself
C: Do you like ketchup?
C: How about pickles?
C: .
S: 250 Message accepted for delivery
C: QUIT
S: 221 hamburger.edu closing connection

Telnet to the Mail Server

telnet <mail server> 25

Mail access protocols



- ❑ SMTP: delivery/storage to receiver's server
- ❑ Mail access protocol: retrieval from server
 - ❖ **POP**: Post Office Protocol [RFC 1939]
 - Authorization (agent <--> server) and download
 - ❖ **IMAP**: Internet Mail Access Protocol [RFC 1730]
 - More features (more complex)
 - Manipulation of stored msgs on server
 - ❖ **HTTP**: gmail, Hotmail, Yahoo! Mail, etc.

POP3 (more) and IMAP

POP3

- ❑ “Download-and-delete” mode: Bob cannot re-read e-mail if he changes client
- ❑ “Download-and-keep” mode: copies of messages are saved on different clients
- ❑ **Stateless** across sessions

IMAP

- ❑ Keep all messages in one place: the server
- ❑ Allows user to organize messages in folders
- ❑ **Stateful** across sessions:
 - ❖ Names of folders and mappings between message IDs and folder name

DNS: Domain Name System

Map between IP addresses and name

- ❑ **IP address:** used for addressing datagrams
- ❑ **Name:** used by humans
 - ❖ e.g., www.yahoo.com

DNS services

- ❑ Hostname to IP address **translation**
- ❑ **Aliasing** (host, mail server)
 - ❖ Canonical, alias names
- ❑ **Load distribution**
 - ❖ Set of IP addresses for one canonical name

Accessing hotmail.com

Local DNS server

No. -	Time	Source	Destination	Protocol	Info
3	1.134841	192.168.1.75	64.59.144.16	DNS	Standard query AAAA login.passport.net
4	1.145866	64.59.144.16	192.168.1.75	DNS	Standard query response CNAME login.passport.net.nsadc.net
5	1.146336	192.168.1.75	64.59.144.16	DNS	Standard query A login.passport.net
6	1.153746	64.59.144.16	192.168.1.75	DNS	Standard query response CNAME login.passport.net.nsadc.net A 65.54.179.196 A 65.54.183.195
7	1.154236	192.168.1.75	65.54.179.196	TCP	45658 > http [SYN] Seq=0 Ack=0 Win=5840 Len=0 MSS=1460 TSV=886259370 TSER=0 WS=2
8	1.189749	65.54.179.196	192.168.1.75	TCP	http > 45658 [SYN, ACK] Seq=0 Ack=1 Win=16384 Len=0 MSS=1460 WS=0 TSV=0 TSER=0
9	1.189798	192.168.1.75	65.54.179.196	TCP	45658 > http [ACK] Seq=1 Ack=1 Win=5840 Len=0 TSV=886259405 TSER=0
10	1.190142	192.168.1.75	65.54.179.196	HTTP	GET /uilogin.srf?lc=1033&id=2 HTTP/1.1
11	1.220375	65.54.179.196	192.168.1.75	TCP	Previous segment lost! http > 45658 [FIN, ACK] Seq=474 Ack=438 Win=65008 Len=0

Frame 6 (149 bytes on wire, 149 bytes captured)
Ethernet II, Src: 00:16:b6:18:b1:37, Dst: 00:0d:88:36:c0:f1
Internet Protocol, Src Addr: 64.59.144.16 (64.59.144.16), Dst Addr: 192.168.1.75 (192.168.1.75)
User Datagram Protocol, Src Port: domain (53), Dst Port: 32918 (32918)
Domain Name System (response)
Transaction ID: 0x2640
Flags: 0x8180 (Standard query response, No error)
Questions: 1
Answer RRs: 3
Authority RRs: 0
Additional RRs: 0
Queries
Answers
login.passport.net: type CNAME, class IN, cname login.passport.net.nsadc.net
Name: login.passport.net
Type: CNAME (Canonical name for an alias)
Class: IN (0x0001)
Time to live: 7 minutes, 24 seconds
Data length: 27
Primary name: login.passport.net.nsadc.net
login.passport.net.nsadc.net: type A, class IN, addr 65.54.179.196
Name: login.passport.net.nsadc.net
Type: A (Host address)
Class: IN (0x0001)
Time to live: 7 minutes, 24 seconds
Data length: 4
Addr: 65.54.179.196
login.passport.net.nsadc.net: type A, class IN, addr 65.54.183.195
Name: login.passport.net.nsadc.net

DNS

Handshake

Http

Canonical hostname

IP address(es)
For hotmail

Returned in different order
each time (load distribution)

Local DNS (Default Name) Server

- ❑ Host's query is sent to its local DNS server
- ❑ Acts as proxy, forwards query into DNS server hierarchy

Unix

> dig nameserver

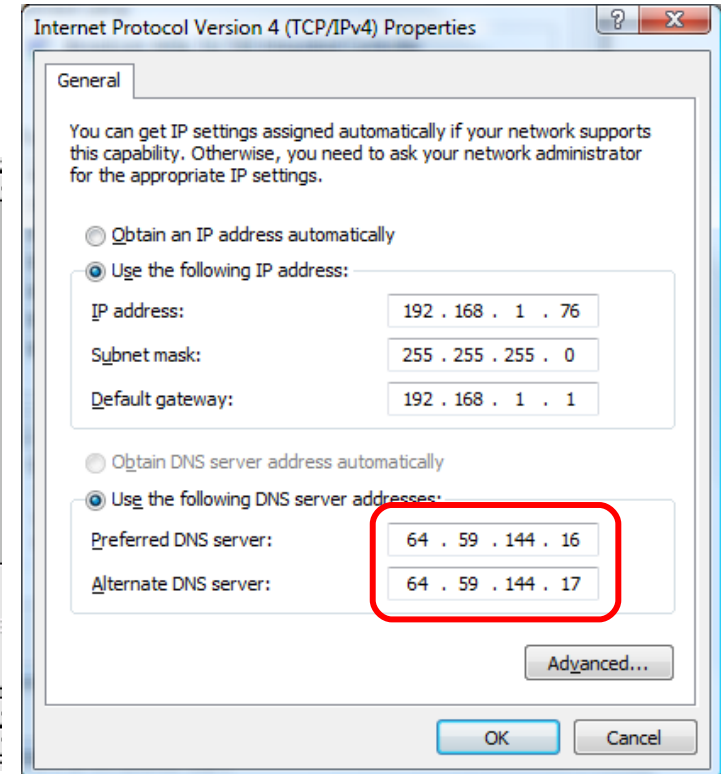
```
; <<>> DiG 9.3.1 <<>> nameserver
;; global options: printcmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NXDOMAIN, id: 28127
;; flags: qr rd ra; QUERY: 1, ANSWER: 0, AUTHORITY: 1, ADDITIONAL: 0

;; QUESTION SECTION:
;nameserver.                IN      A

;; AUTHORITY SECTION:
.                2628    IN      SOA      A.ROOT-SERVERS.NET. NST
GN-GRS.COM. 2007090900 1800 900 604800 86400

;; Query time: 22 msec
;; SERVER: 64.59.144.16#53(64.59.144.16)
;; WHEN: Sun Sep  9 09:38:37 2007
;; MSG SIZE rcvd: 103
```


Windows



Wireshark Lab

1. Download Wireshark
2. Run Wireshark
 - ❖ Capture – Option
 - Interface: eth0 (Linux), (Control Panel - Windows)
 - Capture Filter: src <IP> or dst <IP>
 - “Update list of packets in real time” = yes
 - “Automatic scrolling in live capture” = yes
3. Access to a website

Distributed, Hierarchical Database



□ Top-level domain (TLD) servers:

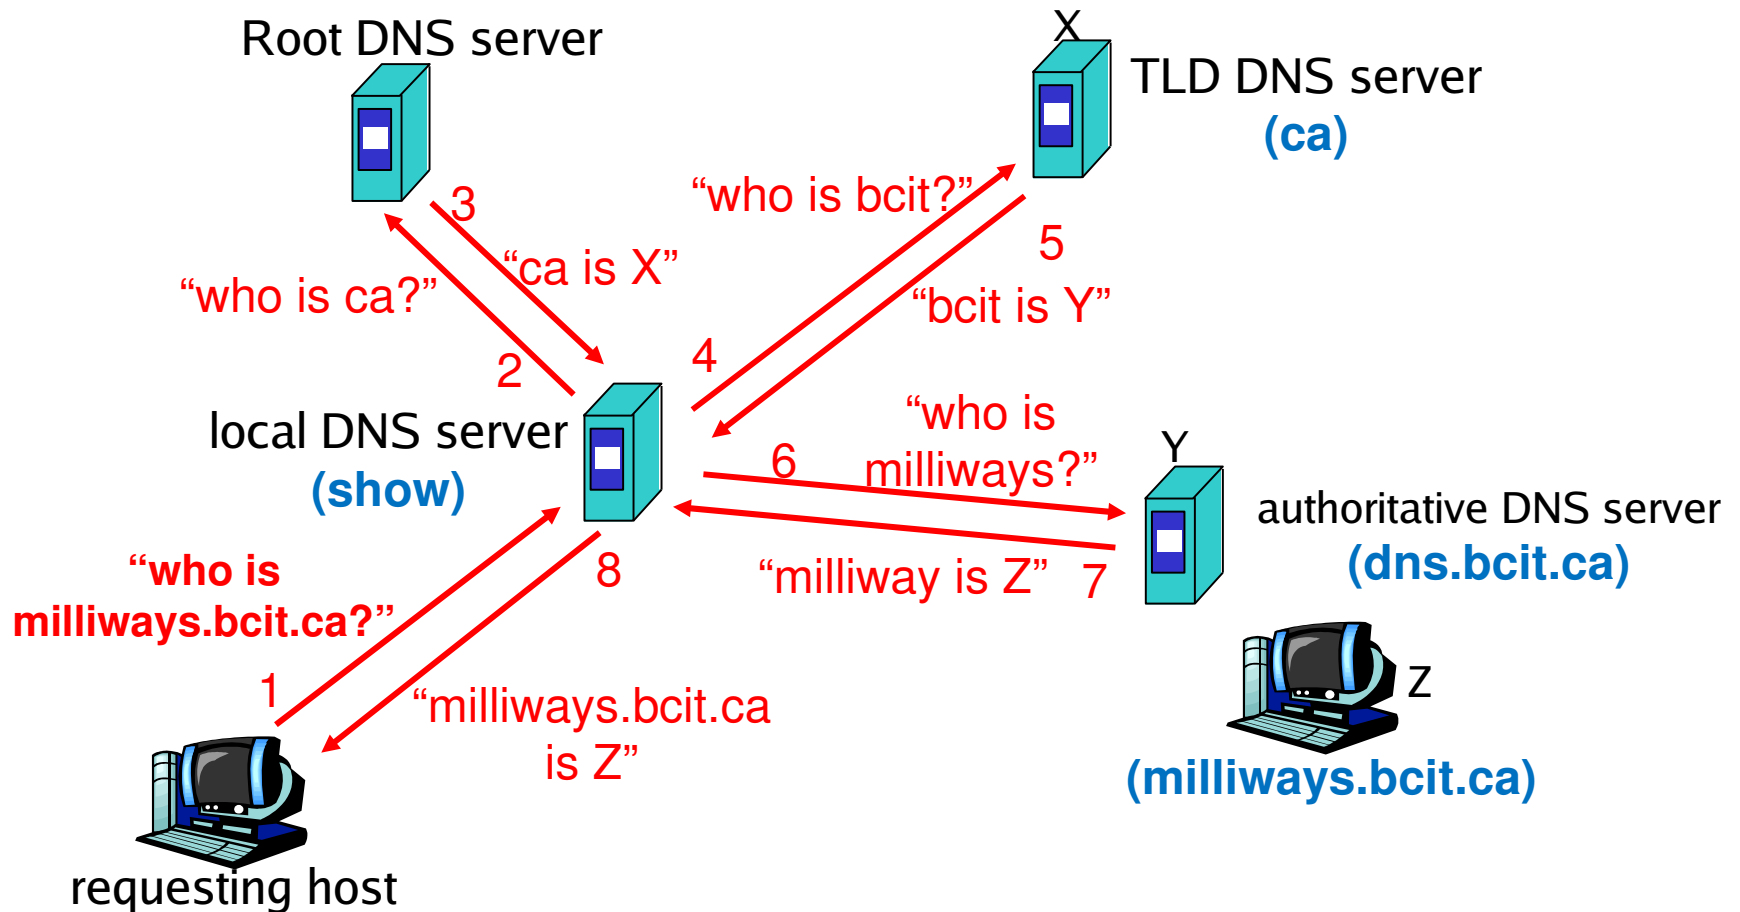
- ❖ Responsible for com, org, net, edu, etc, and all top-level country domains uk, fr, ca, jp.
- ❖ Network Solutions maintains servers for com TLD

□ Authoritative DNS servers:

- ❖ Organization's DNS servers.
- ❖ Can be maintained by organization or service provider

DNS Name Resolution Example

Host at wants IP address for `milliways.bcit.ca`



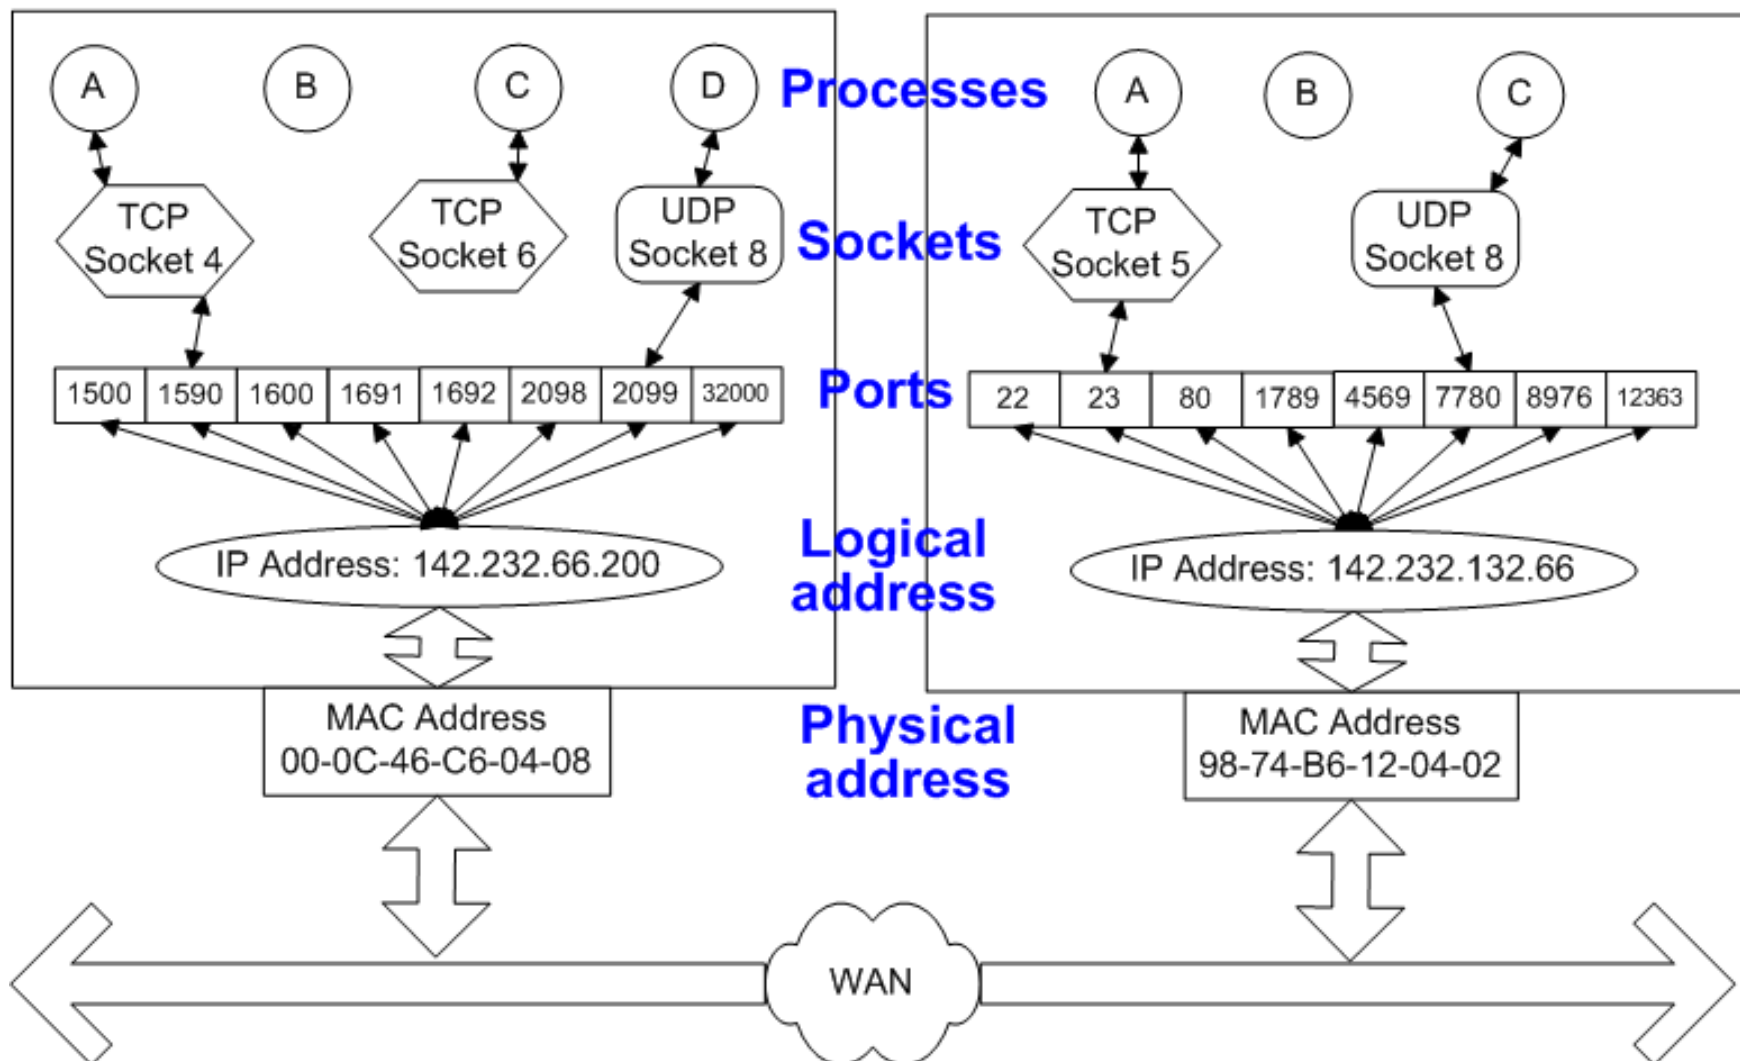
DNS: Caching

- Once (any) name server learns mapping, it *caches* mapping
 - ❖ Cache entries timeout (disappear) after some time
 - ❖ TLD servers typically cached in local name servers
 - => Root name servers not often visited

TCP/IP Client-Server Model

CLIENT

SERVER



Socket

- ❑ Generalization of the **UNIX file access system** designed to incorporate **network protocols**.
- ❑ **Communications channel** between a pair of 'sockets'
- ❑ A socket is defined by a group of four integers:
 - ❖ **Remote host address**
 - ❖ **Remote host port number**
 - ❖ **Local host address**
 - ❖ **Local host port number**

Create Socket

int socket(family, type, protocol)

- ❑ **family (int)**: protocol family
(e.g. **PF_INET**, **PF_UNIX**)
- ❑ **type (int)**: abstract type of the communication desired
(e.g. **SOCK_STREAM**, **SOCK_DGRAM**)
- ❑ **protocol (int)**: specific protocol desired
(e.g. **TCP** or **UDP**)
- ❑ Returns an **socket file descriptor**

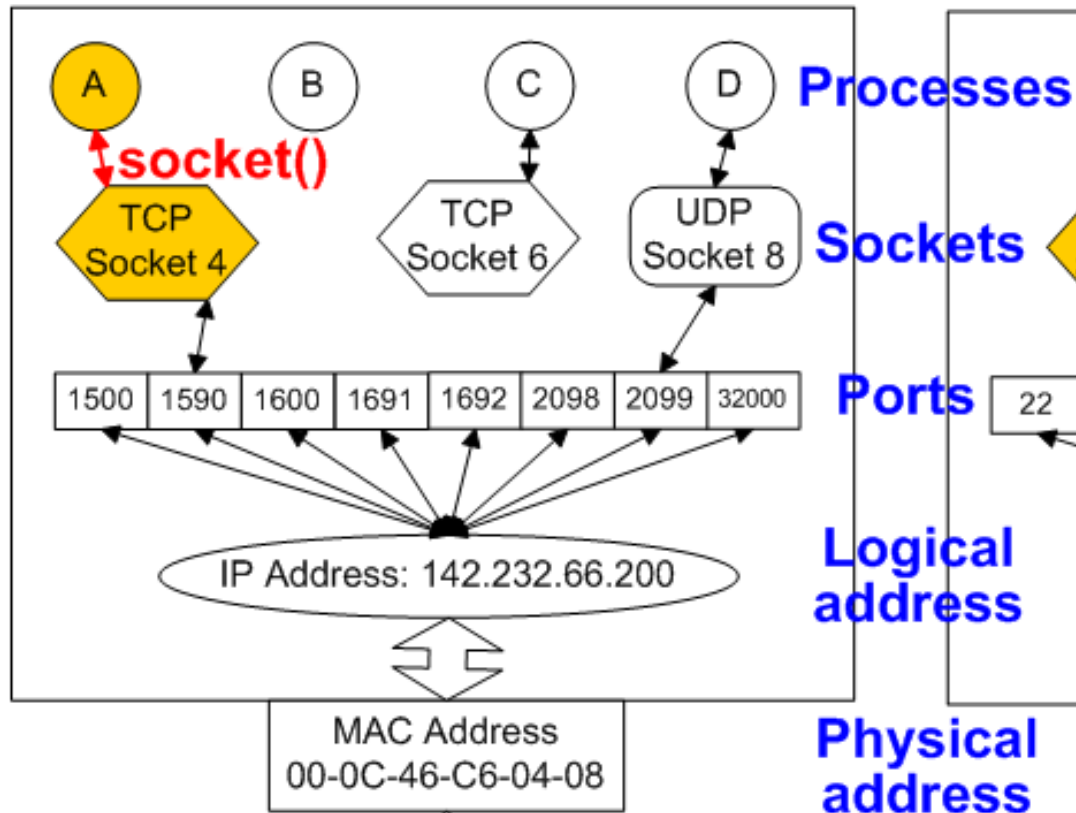
Bind - server

Bind socket to an address

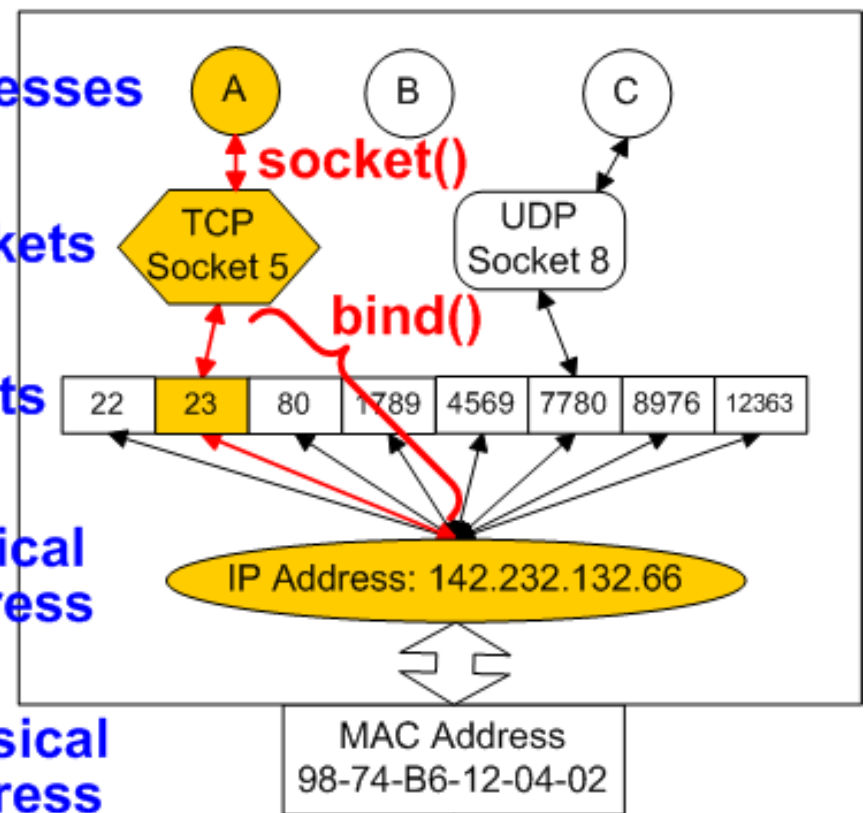
int bind(socket, sockaddr, sockaddrlen)

- ❑ **socket (int)**: socket descriptor
- ❑ **sockaddr (struct sockaddr *)**: **pointer** to the **address** to which the socket should be bound
- ❑ **sockaddrlen (socklen_t)**: the **size** of the **address**
- ❑ Returns -1 and sets **errno** for error

CLIENT



SERVER



listen() - server

Marks socket as 'listening' and sets the maximum number of listen queue

int listen(sock, backlog)

- ❑ **sock (int)**: socket
- ❑ **backlog (unsigned int)**: max length of the queue of unprocessed connection requests

Addressing

- ❑ **common framework** for all addresses to support multiple protocol families
- ❑ In the **Internet family**, transport addresses are **6 bytes** long:
 - ❖ **4 bytes** for the **Internet address**
 - ❖ **2 bytes** for the **port number**.
- ❑ **<netinet/in.h>** defines the appropriate structures to use for the Internet family

```
struct sockaddr_in
```

```
{
```

```
    short sin_family;
```

```
    unsigned short sin_port;
```

```
    struct in_addr sin_addr;
```

```
    unsigned char sin_zero[8]; /* unused */
```

```
}
```

sin_family: address family

sin_port: 16-bit port number

sin_addr: 32-bit Internet address

```
struct in_addr
```

```
{
```

```
    u_long s_addr; /* 32-bit IP Address */
```

```
}
```

Get Server Information

gethostbyname(host_name)

- ❑ Maps the request into a **DNS** query that it sends to the **resolver** running on the local machine.
- ❑ Returns a **pointer** to a **hostent** structure that contains the requested addresses.
- ❑ The related function **gethostbyaddr()** takes an **Internet address** as an argument instead.


```
struct hostent
{
    char *h_name;
    char **h_alias;
    int h_addrtype;
    int h_length;
    char **h_addr_list;
};
```

h_name: the official, fully qualified **domain name**

h_alias: a null-terminated list of alternate names
(aliases)

h_addrtype: **protocol family** the address belongs to

h_length: **length** of the address

h_addr_list: a null-terminated **array of addresses**

Address Manipulation Functions

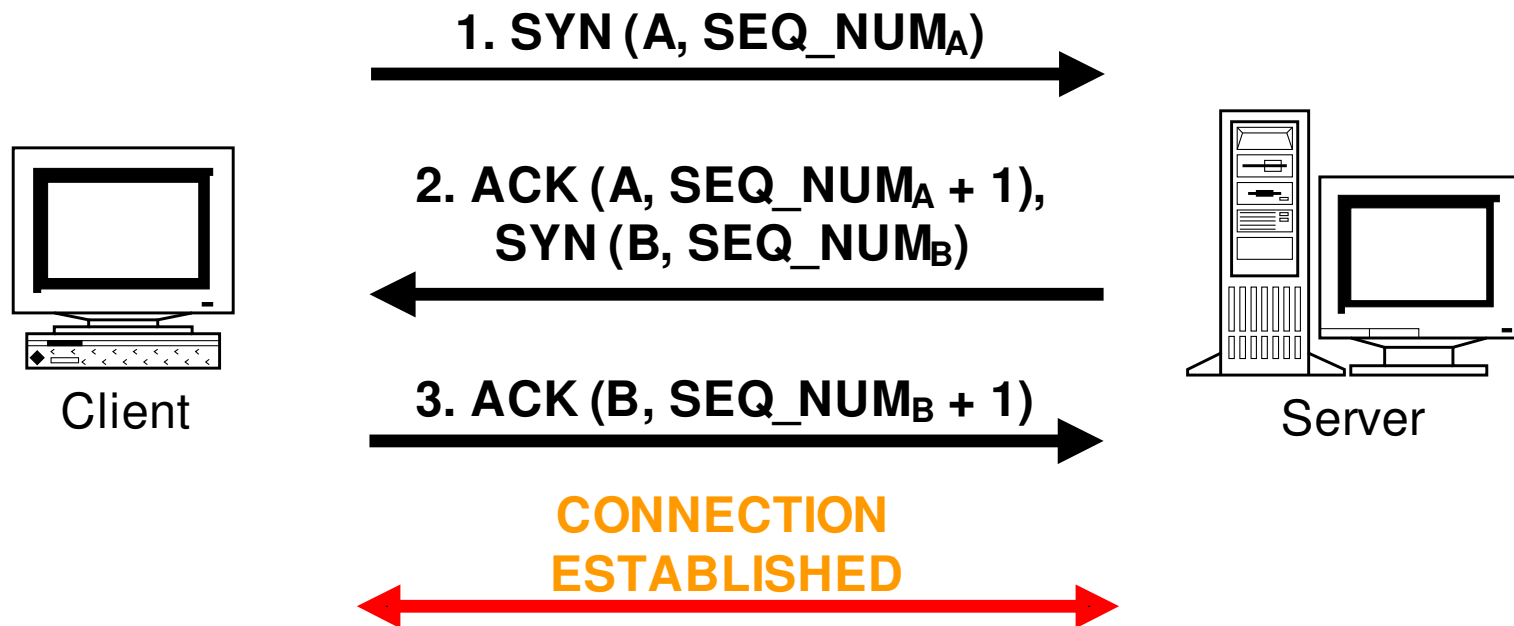
inet_aton() / inet_network()

- ❑ Takes strings representing Internet addresses in **dotted notation**
- ❑ Returns numbers suitable for use in **sockaddr_in** structures.

inet_ntoa()

- ❑ Takes a 32-bit IP address in **network byte order** Internet address
- ❑ Returns the corresponding address in dotted-decimal notation.

Establish a Connection to Server



TCP 3-Way Handshake

Send Connect Request (SYN)

**connect(int sock,
 struct sockaddr *name,
 int namelen)**

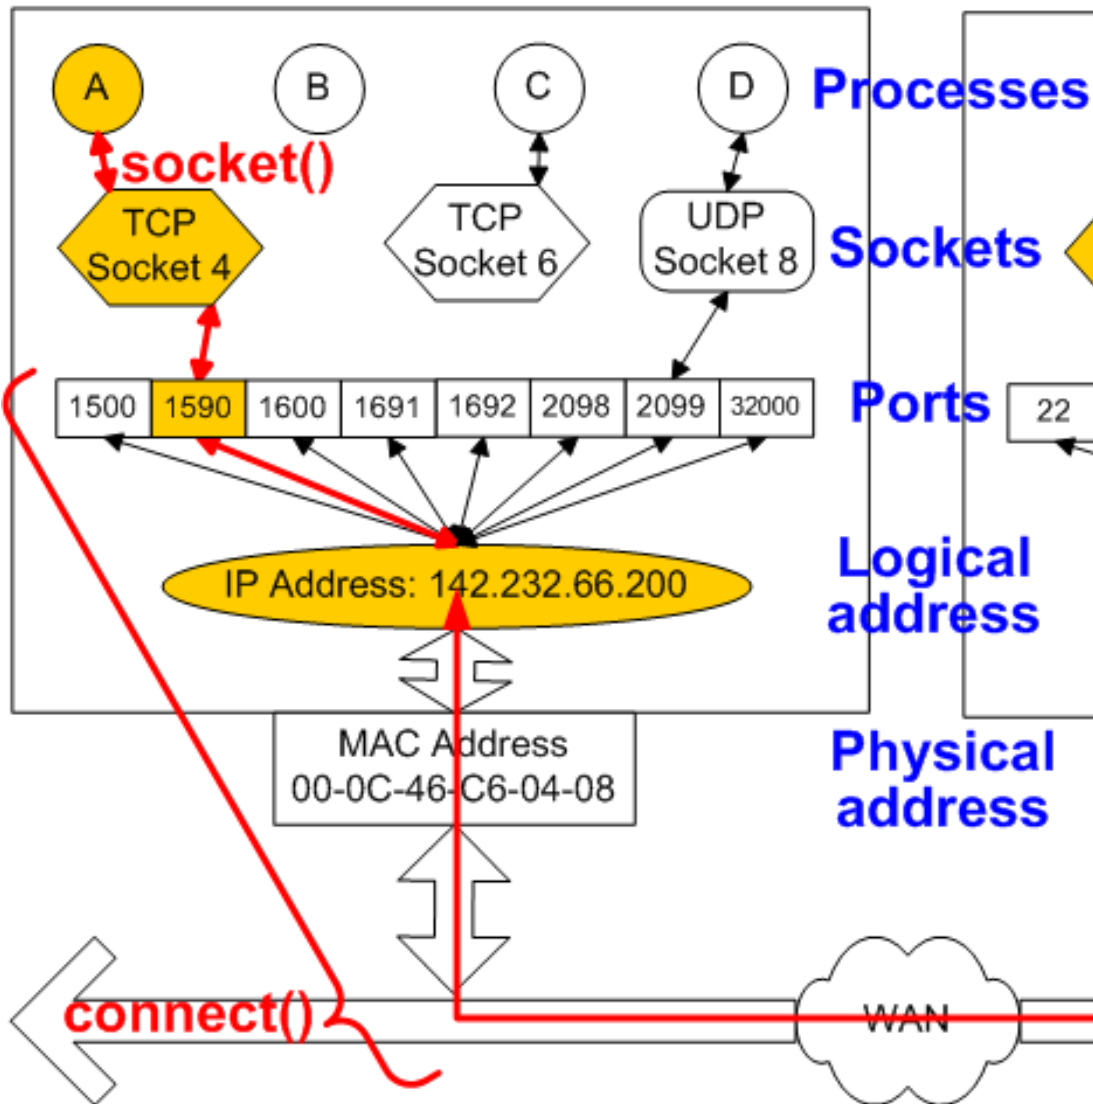
- ❑ Establishes communication with a remote entity
 - ❖ initiates a **three-way handshake** for TCP connection
 - ❖ normally only used for **SOCK_STREAM** sockets.
- ❑ **sock**: client socket
- ❑ **name**: points to an area of memory containing the address information for the remote entity
- ❑ **namelen**: length of 'name'
- ❑ Returns -1 and sets **errno** for error

Accept Connect Request (SYN/ACK)

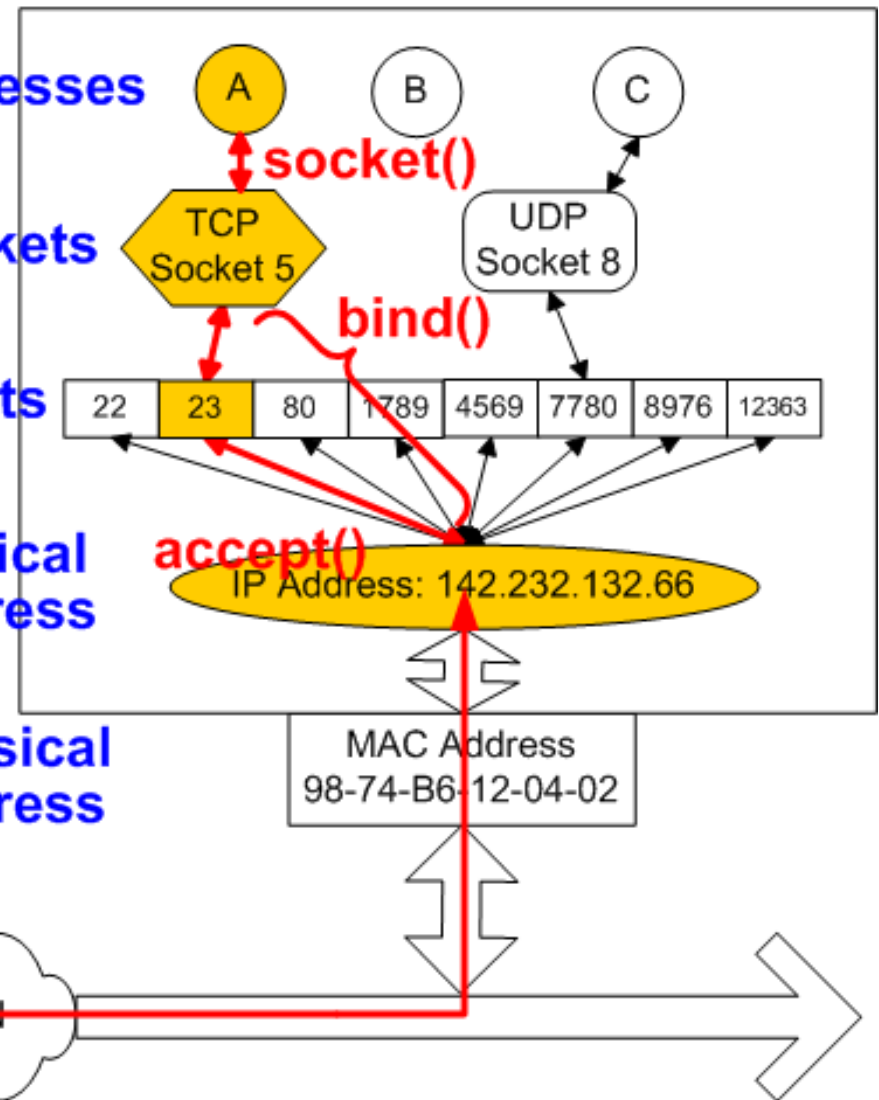
**accept (int sock,
 struct sockaddr *addr,
 int *addrlen)**

- ❑ **sock**: base socket descriptor
- ❑ **addr**: points to a struct sockaddr of the remote (client) system.
- ❑ **addrlen**: size of addr
- ❑ Returns a new socket descriptor (int) that will be used for all subsequent communication with the remote host.

CLIENT



SERVER



Send/Receive Data

send(sock, message, length, flags)
recv(sock, message, length, flags)

- ❑ **sock**: socket descriptor
- ❑ **message**: a pointer to the data to be sent/received
- ❑ **length**: data length.
- ❑ **flags**: to send **flags** to the **underlying protocol**.
 - ❖ flags may be formed by ORing MSG_OOB and MSG_DONTROUTE.
- ❑ send() and recv() may only be used with **SOCK_STREAM** type sockets.
- ❑ Returns the number of bytes transmitted or -1 for error

Send/Receive Data - UDP

**sendto(sock, message, length, flags, dest,
destlength)**

**recvfrom(sock, message, length, flags, from,
fromlength)**

- The two additional arguments are **pointers** to a **sockaddr** structure and its **length**.

Close Connection

shutdown (socket, how)

to shut down all or part of a full-duplex connection

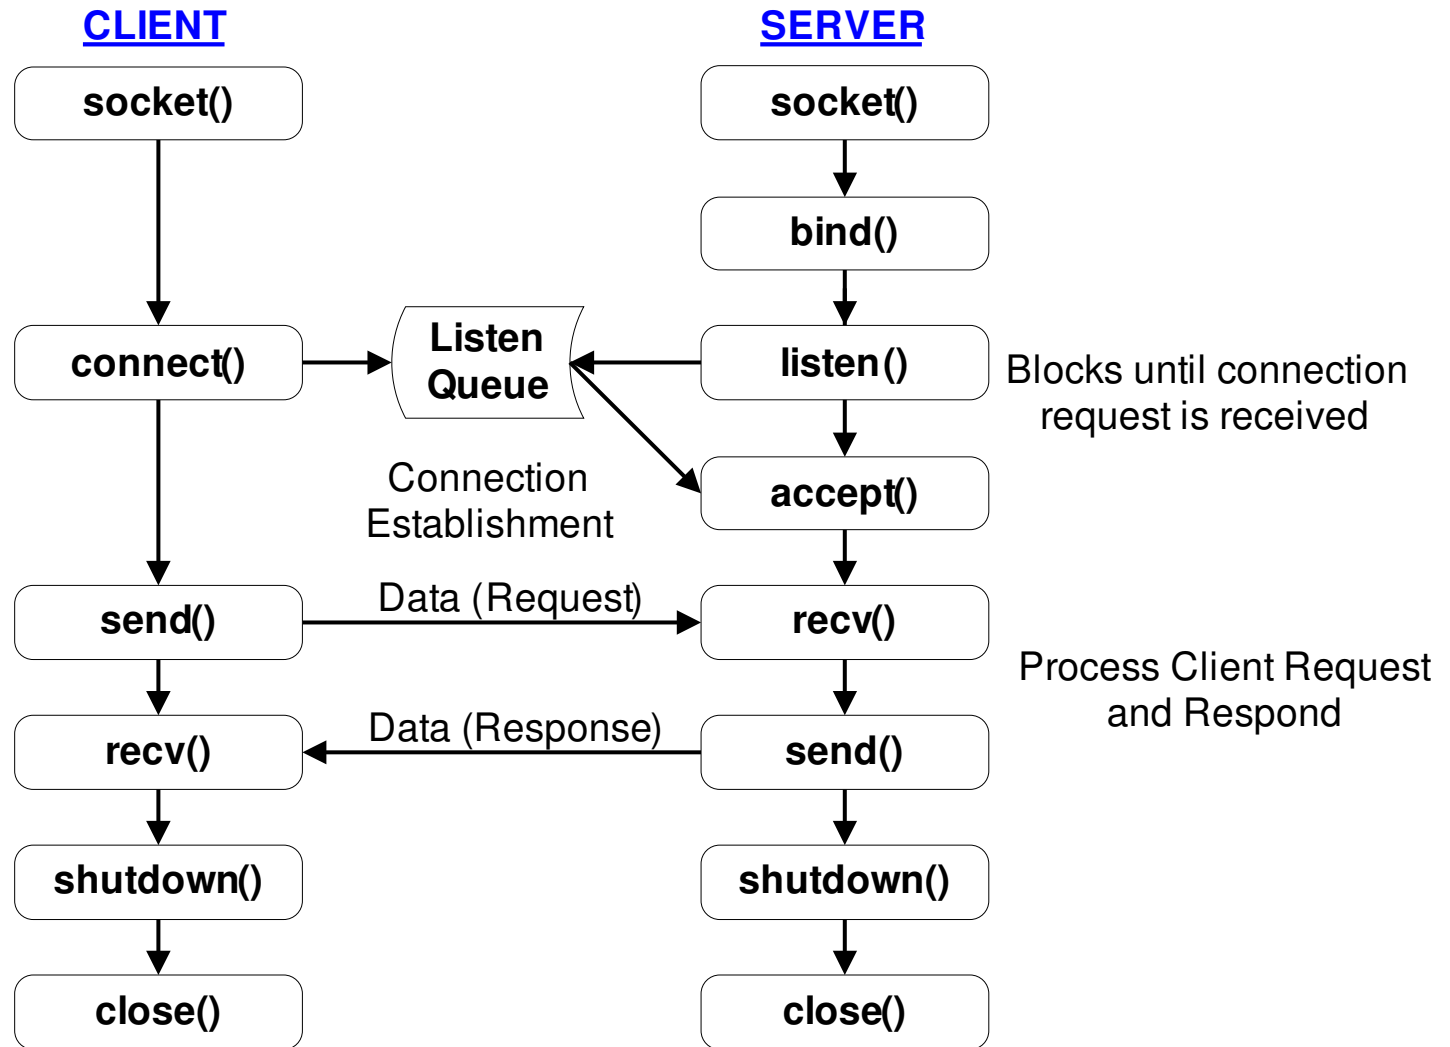
□ how: **direction** of the connection to be shutdown:

- ❖ 0 - additional receives are shutdown.
- ❖ 1 - additional sends are disallowed.
- ❖ 2 - additional sends and receives will be disabled.

close(socket)

to close a socket

TCP/IP Client Server Model



UDP Client Server Model

